

StatML LAB-1

Import numpy and pandas for operations

```
import numpy as np
import pandas as pd
```

'np.array()' function produces an array as shown

```
lst = [1,2,3] array1 =
np.array(lst) array1
```

```
array([1, 2, 3])
```

And as for the given functions 'np' produces a set of zeros, ones and range of numbers between certain numbers.

0	abc	12	70
1	xyz	13	80

```
2      hij      14      90 print("A series of zeroes:",np.zeros(7)) print("A
series of ones:",np.ones(9)) print("A series of
numbers:",np.arange(5,16)) print("Numbers spaced apart by
2:",np.arange(0,11,2)) print("Numbers spaced apart by
float:",np.arange(0,11,2.5)) print("Every 5th number from 30 in
reverse order: ",np.arange(30,-1,print("11 linearly spaced numbers
between 1 and 5: ",np.linspace(1,5,
A series of zeroes: [0. 0. 0. 0. 0. 0. 0.]
A series of ones: [1. 1. 1. 1. 1. 1. 1. 1. 1.]
A series of numbers: [ 5  6  7  8  9 10 11 12 13 14 15]
Numbers spaced apart by 2: [ 0  2  4  6  8 10]
Numbers spaced apart by float: [ 0.   2.5  5.   7.5 10. ]
Every 5th number from 30 in reverse order: [30 25 20 15 10  5  0]
11 linearly spaced numbers between 1 and 5: [1.  1.4 1.8 2.2 2.6 3.   3.4 3.8 4.2 4.6 5. ]
```

'pd.read_csv' function reads a CSV le and syntax as shown.

```
dataframe = pd.read_csv("wine.csv")
dataframe.head()
```

	Wine	Alcohol	Malic.acid	Ash	Ac1	Mg	Phenols	Flavanoids	Nonflavanoid.phenols	Proanth	Color.int	Hue	OD	Proline
0	1	14.23	1.71	2.43	15.6	127	2.80	3.06	0.28	2.29	5.64	1.04	3.92	1065
1	1	13.20	1.78	2.14	11.2	100	2.65	2.76	0.26	1.28	4.38	1.05	3.40	1050
2	1	13.16	2.36	2.67	18.6	101	2.80	3.24	0.30	2.81	5.68	1.03	3.17	1185
3	1	14.37	1.95	2.50	16.8	113	3.85	3.49	0.24	2.18	7.80	0.86	3.45	1480
4	1	13.24	2.59	2.87	21.0	118	2.80	2.69	0.39	1.82	4.32	1.04	2.93	735

```
datacsv = pd.read_csv("createcsv.csv")
datacsv
```

name Location Marks
'pd.read_txt' function reads a text le and syntax as shown.

```
datatext = pd.read_table("pandatext.txt")
datatext
```

	Name	Height	Weight	Hometown
0	0 Ashley	155	140	Palo Alto
1	1 Robin	145	122	Fremont
2	2 Priyanka	152	131	Santa Clara
3	3 Youngchul	167	148	Cupertino
4	4 Aziz	161	139	San Francisco
5	5 Zoey	181	190	Hayward

'pd.read_xlsx' function reads a excel le and syntax as shown.

```
datax1 = pd.read_excel("Height_weight.xlsx")
datax1
```

	Name	Height	Weight
0	Ashton	155	135
1	Kate	125	140
2	Bruce	178	210
3	Tom	181	165
4	Bill	165	180


```
read_from_html = pd.read_html("")
```

Now let's import 'matplotlib.pyplot' for graph operations. And given are the examples of this library like scatter, hist etc.



```
import matplotlib.pyplot as plt
```

```
people = ['Ann', 'Brandon', 'Chen', 'David', 'Emily', 'Farook',  
'Gagan', 'Hamish', 'Imran', 'Julio', 'Katherine', 'Lily'] age =  
[21, 12, 32, 45, 37, 18, 28, 52, 5, 40, 48, 15] weight =  
[55, 35, 77, 68, 70, 60, 72, 69, 18, 65, 82, 48] height =  
[160, 135, 170, 165, 173, 168, 175, 159, 105, 171, 155, 158]
```

```
plt.scatter(age,weight)  
plt.show()
```



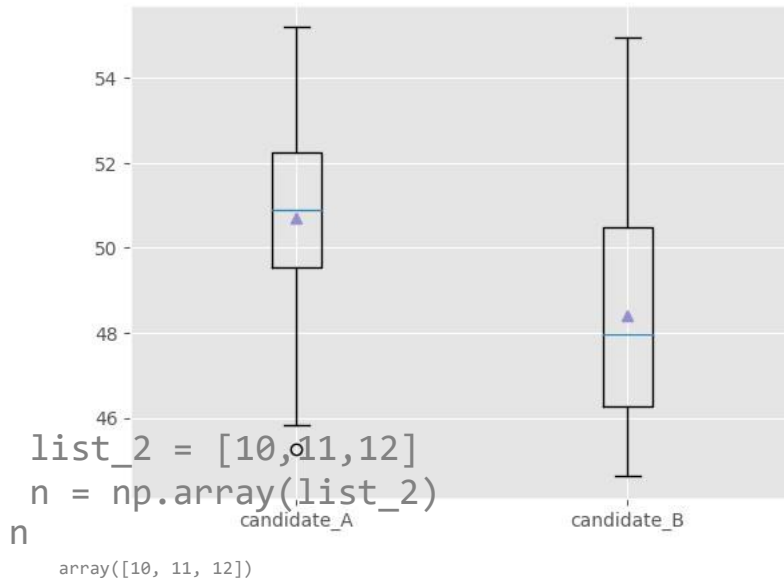
```
plt.hist(age)  
plt.show()
```



```

days = np.arange(1,31) candidate_A =
50+days*0.07+2*np.random.randn(30) candidate_B = 50-
days*0.1+3*np.random.randn(30) plt.style.use("ggplot")
plt.boxplot(x=[candidate_A,candidate_B],showmeans = True)
plt.grid(True) plt.xticks([1,2],["candidate_A","candidate_B"])
plt.show()

```



```

list_2 = [10,11,12]
n = np.array(list_2)
n
array([10, 11, 12])

```

```

print(f"Sum of arrays {lst} and {list_2} is {array1 + n}")
Sum of arrays [1, 2, 3] and [10, 11, 12] is [11 13 15]

```

```

#Operations on numpy arrays print(f"Multiplication of Numpy arrays
{array1 * n}") print(f"Subtraction of Numpy arrays {array1 - n}")
print(f"Division of Numpy arrays {array1 / n}")

```

```

Multiplication of Numpy arrays [10 22 36]
Subtraction of Numpy arrays [-9 -9 -9]
Division of Numpy arrays [0.1      0.18181818 0.25      ]

```

StatML LAB-2

Import required libraries

```
import numpy as np
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D
from scipy.stats import multivariate_normal
```

First, you define a 4x4 covariance matrix named `covariance`. This matrix represents the covariances between four variables. The covariance matrix contains the variances of each variable on the diagonal and the covariances between variables off the diagonal.

You then use `np.linalg.inv(covariance)` to calculate the inverse of the covariance matrix. The `np.linalg.inv` function is a NumPy function that computes the matrix inverse.

Finally, you print the result, which is the precision matrix.

```
covariance = np.array([[0.14, -0.3, 0.0, 0.2],
                       [-0.3, 1.16, 0.2, -0.8],
                       [0.0, 0.2, 1.0, 1.0],
                       [0.2, -0.8, 1.0, 2.0]])
precision = np.linalg.inv(covariance)
print(precision)
```

```
[[ 60.    50.   -48.    38.]
 [ 50.    50.   -50.    40.]
 [-48.   -50.   52.4  -41.4]
 [ 38.    40.  -41.4   33.4]]
```

`generate_pair()` is a function that generates a pair of random values following a bivariate (two-dimensional) normal distribution. The parameters passed to `np.random.multivariate_normal` specify the mean vector and the covariance matrix for the distribution:

The mean vector `[0.8, 0.8]` indicates that the distribution is centered around the point `(0.8, 0.8)`. This means that on average, the generated pairs will tend to be close to this point.

The covariance matrix `[[0.1, -0.1], [-0.1, 0.12]]` specifies how the two variables (in this case, `x` and `y`) are related. The values on the diagonal (`0.1` and `0.12`) represent the variances of the two variables, while the off-diagonal values (`-0.1`) represent the covariance between the variables. This matrix determines how spread out or correlated the generated pairs will be. Positive covariance values indicate that the variables tend to increase together, while negative values indicate that they tend to move in opposite directions.

When you call `generate_pair()`, it returns a pair of random values drawn from the specified bivariate normal distribution, and the values are stored in the variable `mu_t`.

Finally, you print the `mu_t` variable, which displays the pair of random values generated by the function.

```
def generate_pair():
    return np.random.multivariate_normal([0.8, 0.8], [[0.1, -0.1], [-0.1, 0.12]])
```

```
mu_t = generate_pair()
print(mu_t)
```

```
[0.93474884 0.80350139]
```

```
x, y = np.mgrid[-0.25:2.25:.01, -1:2:.01]
pos = np.empty(x.shape + (2,))
pos[:, :, 0] = x
pos[:, :, 1] = y
```

https://colab.research.google.com/drive/1mWcNE9QqgOBcevQYs-EO_h9WqychaQ5h#scrollTo=AH1nZDtF1Pf2&printMode=true

```
mu_p = np.zeros((1000, 2))
mu_p[:, 0] = 0.8
mu_p[:, 1] = 0.8
```

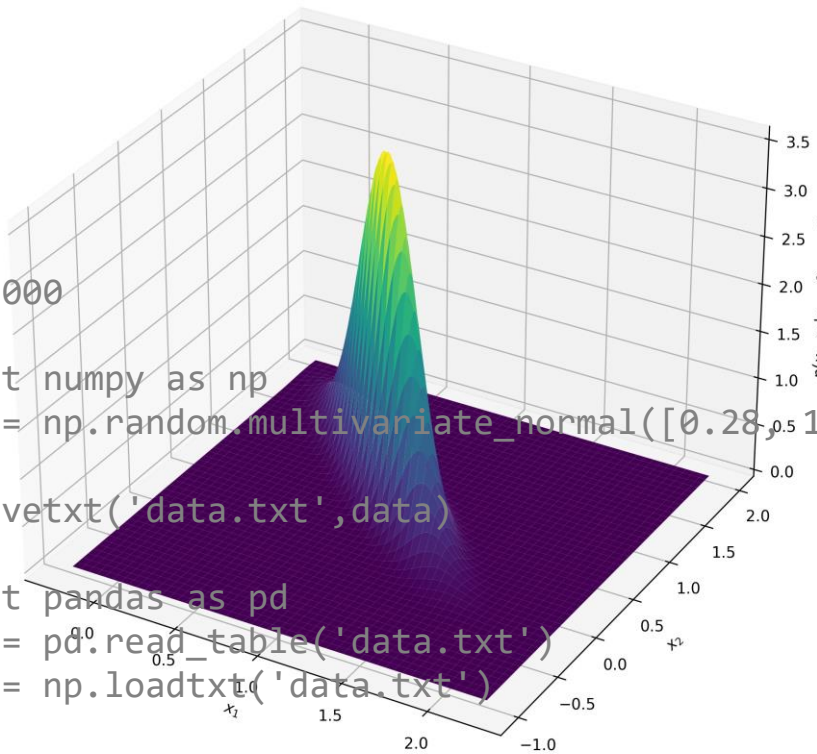
```
mu_p = [0.8, 0.8] cov_p = [[0.1, -0.1], [-0.1, 0.12]] z =  
multivariate_normal(mu_p, cov_p).pdf(pos)
```

```
fig = plt.figure(figsize=(10, 10), dpi=300)  
ax = fig.add_subplot(projection='3d')  
ax.plot_surface(x, y, z, cmap=plt.cm.viridis)  
plt.xlabel('$x_1$') plt.ylabel('$x_2$')  
ax.set_zlabel('$p(x_1, x_2 \mid x_3=0, x_4=0)$')  
plt.savefig('cond_mvg.png', bbox_inches='tight', dpi=300) plt.show()
```

N = 1000

```
import numpy as np
data = np.random.multivariate_normal([0.28, 1.18], [[2.0, 0.8], [0.8, 2.0]])
np.savetxt('data.txt', data)

import pandas as pd
data = pd.read_table('data.txt')
data = np.loadtxt('data.txt')
```



data

data

```
array([[ -1.97416035,  0.18085742],
 [ 1.89414551, -0.14292725],
 [ 2.44416774, -1.48606541],
 ...,
 [ 2.2292759 ,  0.26630475],
 [-1.50839992, -0.01755483],
 [-1.63801017,  2.59365018]])
```

```
mu_ml = data.mean(axis=0) x = data - mu_ml
cov_ml = np.dot(x.T, x) / N
cov_ml_unbiased = np.dot(x.T, x) / (N - 1)
print(mu_ml) print(cov_ml)
print(cov_ml_unbiased)
```

```
[0.32950325  1.24389006]
[[2.02971382  0.8557863 ]
 [0.8557863  4.23663802]]
[[2.03174557  0.85664294]
 [0.85664294  4.24087889]]
```



```
def seq_ml(data):      mus = [np.array([[0],
[0]])]      for i in range(N):      x_n =
data[i].reshape(2, 1)      mu_n = mus[-1] +
(x_n-mus[-1]) / (i + 1)      mus.append(mu_n)
return mus
```

```
mus_ml = seq_ml(data) print(mus_ml[-
1])
[[0.32950325]
[1.24389006]]
```

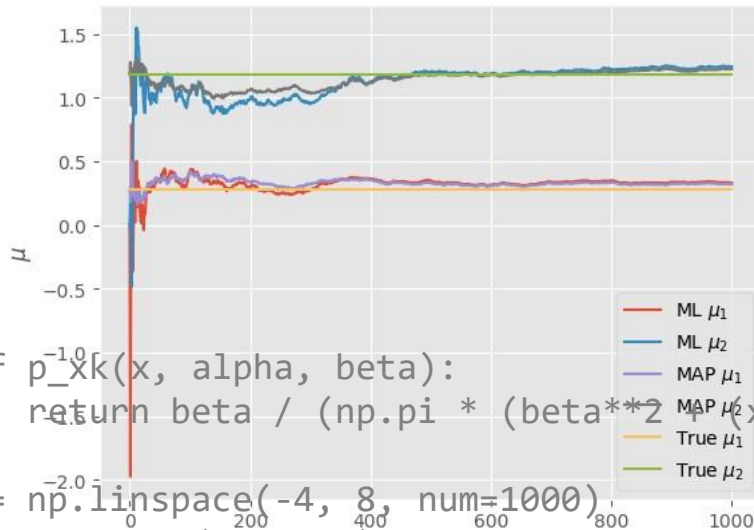
```
mu_p = np.array([[0.28], [1.18]]) cov_p = np.array([[0.1, -0.1], [-
0.1, 0.12]]) cov_t = np.array([[2.0, 0.8], [0.8, 4.0]])
```

```
def seq_map(data, mu_p, cov_p, cov_t):
mus, covs = [mu_p], [cov_p]      for x
in data:      x_n = x.reshape(2, 1)
cov_n = np.linalg.inv(np.linalg.inv(covs[-1]) + np.linalg.inv
mu_n = cov_n.dot(np.linalg.inv(cov_t).dot(x_n) + np.linalg.in
mus.append(mu_n)      covs.append(cov_n)      return mus, covs
```

```
mus_map, covs_map = seq_map(data, mu_p, cov_p, cov_t) print(mus_map[-
1])
[[0.31735853]
[1.22405254]]
```

```
X = np.arange(N+1)
mus1_ml = [mu[0] for mu in mus_ml]
mus2_ml = [mu[1] for mu in mus_ml]
mus1_map = [mu[0] for mu in mus_map]
mus2_map = [mu[1] for mu in mus_map]
mus1_t = [0.28] * (N+1) mus2_t =
[1.18] * (N+1)
plt.style.use('ggplot')
plt.plot(X, mus1_ml, label='ML $\mu_1$')
plt.plot(X, mus2_ml, label='ML $\mu_2$')
plt.plot(X, mus1_map, label='MAP $\mu_1$')
plt.plot(X, mus2_map, label='MAP $\mu_2$')
plt.plot(X, mus1_t, label='True $\mu_1$')
plt.plot(X, mus2_t, label='True $\mu_2$')
plt.xlabel('$n$-th data point')
plt.ylabel('$\mu$') plt.legend(loc=4)
```


```
plt.savefig('seq_learning.png', bbox_inches='tight', dpi=300)
plt.show()
```



```
def p_xk(x, alpha, beta):
    return beta / (np.pi * (beta**2 + (x-alpha)**2))
```


```
x = np.linspace(-4, 8, num=1000)
probs = p_xk(x, 2, 1)
```

```
plt.plot(x, probs)
plt.xlabel('$x_k$')
plt.ylabel(r'$p(x_k | \alpha, \beta)$')
plt.savefig('prob_xk.png',
bbox_inches='tight', dpi=300)
plt.show()
```



```
def p_a(x, alpha, beta):
    return np.product(beta / (np.pi *
beta**2 + (x-alpha)**2))
```

```
D = np.array([4.8, -2.7, 2.2, 1.1, 0.8, -7.3])
alphas = np.linspace(-5, 5, num=1000) beta = 1
likelihoods = [p_a(D, alpha, beta) for alpha
in alphas] plt.plot(alphas, likelihoods)
plt.xlabel(r'\alpha$') plt.ylabel(r'$p(a | D,
\beta)$') plt.savefig('prob_a.png',
bbox_inches='tight', dpi=300) plt.show()
```



```
print(D.mean())
print(alphas[np.argmax(likelihoods)])
-0.18333333333333326
1.1761761761761758
```

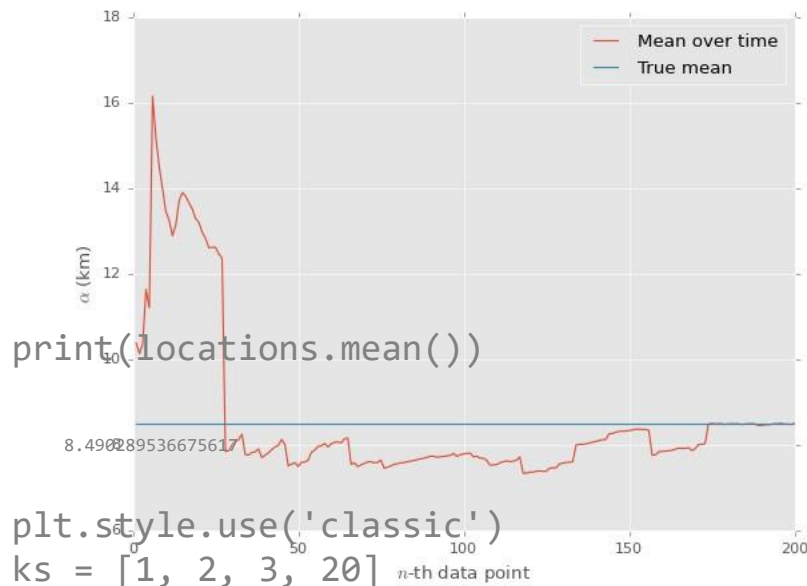
```
alpha_t = np.random.uniform(0, 10)
beta_t = np.random.uniform(1, 2)
print(alpha_t, beta_t)
9.556577442054802 1.3712422301299707
```

```
def location(angle, alpha, beta):
return beta * np.tan(angle) + alpha
```

```
N = 200 angles = np.random.uniform(-np.pi/2, np.pi/2, N) locations =
np.array([location(angle, alpha_t, beta_t) for angle in a
```

```
mus = [locations[:i + 1].mean() for i in range(N)] mean =
[locations.mean()] * (N) X = np.arange(1, N + 1)
plt.style.use('ggplot')
plt.plot(X, mus, label='Mean over time')
plt.plot(X, mean, label='True mean')
plt.xlabel('$n$-th data point')
plt.ylabel(r'\alpha$ (km)')
plt.legend()
```

```
plt.savefig('mean_x.png', bbox_inches='tight', dpi=300)
plt.show()
```



```
plt.style.use('classic')
ks = [1, 2, 3, 20]
alphas, betas = np.mgrid[-10:10:0.04, 0:5:0.04] # alphas, betas =
np.meshgrid(np.linspace(-10, 10, num=500), np.linspace(0, 5, num=500))
x = locations[:,k] # We only have to calculate the constant once
likelihood = k * np.log(betas/np.pi) for loc in x:
likelihood -= np.log(betas**2 + (loc - alphas)**2)
```

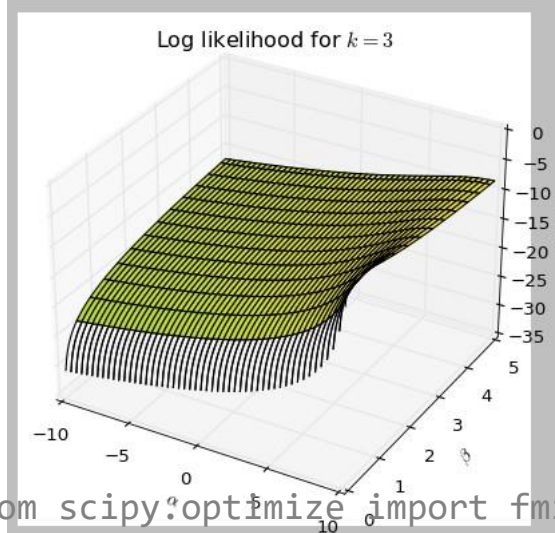
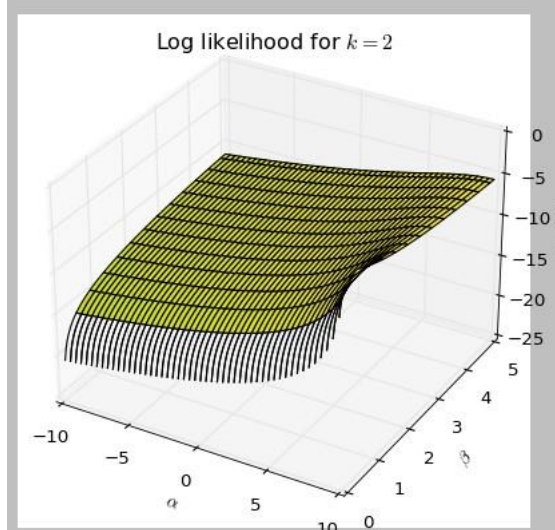

```
fig = plt.figure() ax = fig.add_subplot(projection='3d')
ax.plot_surface(alphas, betas, likelihood, cmap=plt.cm.viridis, v
plt.xlabel(r'$\alpha$') plt.ylabel(r'$\beta$')
ax.set_zlabel('$\ln p(D | \alpha, \beta)$') plt.title('Log
likelihood for $k = \{k\}$.format(k))
plt.savefig('logl_{}.png'.format(k), bbox_inches='tight', dpi=300)
plt.show()
```

<ipython-input-49-f58e94a921d6>:8: RuntimeWarning: divide by zero encountered in log

likelihood = k * np.log(betas/np.pi)

/usr/local/lib/python3.10/dist-packages/mpl_toolkits/mplot3d/proj3d.py:180: RuntimeWarning: invalid value encountered in true_divide
txs, tys, tzs = vecw[0]/w, vecw[1]/w, vecw[2]/w

```
likelihood -= np.log
```



```
from scipy.optimize import fmin
```

```
def log_likelihood(params, locations):
```

```
    alpha, beta = params
```

```
    likelihood = len(locations) * np.log(beta/np.pi)
```

```
    for loc in locations:
```

```
        likelihood -= (beta**2 + (loc -
```

```
        alpha)**2)
```

```
        (beta**2 + (loc -
```

```

def plot_maximize_logl(data, alpha_t, beta_t):
    alphas, betas = [], []
    x = np.arange(len(data))
    for k in x:
        [alpha, beta] = fmin(log_likelihood, (0, 1), args=(data[:k],))
        alphas.append(alpha)
        betas.append(beta)

    plt.style.use('ggplot')
    plt.plot(x, alphas, label=r'$\alpha$')
    plt.plot(x, betas, label=r'$\beta$')
    plt.plot(x, [alpha_t]*len(data), label=r'$\alpha_t$')
    plt.plot(x, [beta_t]*len(data), label=r'$\beta_t$')
    plt.xlabel('$k$')
    plt.ylabel('location (km)')
    plt.legend()
    plt.savefig('plots/min_logl.png', bbox_inches='tight', dpi=300)
    plt.show()
    print(alphas[-1], betas[-1])

plot_maximize_logl(locations, alpha_t, beta_t)

```

```
Optimization terminated successfully.
  Current function value: 0.000000
  Iterations: 10
  Function evaluations: 39
Optimization terminated successfully.
  Current function value: 0.898559
  Iterations: 104
  Function evaluations: 197
Optimization terminated successfully.
  Current function value: 2.181169
  Iterations: 86
  Function evaluations: 162
Optimization terminated successfully.
  Current function value: 7.644875
  Iterations: 85
  Function evaluations: 162
Optimization terminated successfully.
  Current function value: 9.410418
  Iterations: 87
  Function evaluations: 164
Optimization terminated successfully.
  Current function value: 17.834954
  Iterations: 86
  Function evaluations: 162
Optimization terminated successfully.
  Current function value: 19.446122
  Iterations: 94
  Function evaluations: 177
Optimization terminated successfully.
  Current function value: 20.190702
  Iterations: 94
  Function evaluations: 171
Optimization terminated successfully.
  Current function value: 20.627525
  Iterations: 90
  Function evaluations: 172
Optimization terminated successfully.
  Current function value: 22.856079
  Iterations: 77
  Function evaluations: 150
Optimization terminated successfully.
  Current function value: 25.308925
  Iterations: 89
  Function evaluations: 168
<ipython-input-45-6a1108c4c769>:5: RuntimeWarning: invalid value encountered in log
likelihood = len(locations) * np.log(beta/np.pi)
<ipython-input-45-6a1108c4c769>:14: RuntimeWarning: Maximum number of function evaluations has been exceeded.
  [alpha, beta] = fmin(log_likelihood, (0, 1), args=(data[:k],))
Optimization terminated successfully.
  Current function value: 27.014560
  Iterations: 80
  Function evaluations: 150
Optimization terminated successfully.
  Current function value: 32.265839
  Iterations: 85
  Function evaluations: 159
Optimization terminated successfully.
  Current function value: 38.483444
  Iterations: 87
  Function evaluations: 162
Optimization terminated successfully.
  Current function value: 43.508164
  Iterations: 86
  Function evaluations: 163
Optimization terminated successfully.
  Current function value: 46.596925
  Iterations: 98
  Function evaluations: 186
Optimization terminated successfully.
  Current function value: 48.454299
  Iterations: 83
  Function evaluations: 157
Optimization terminated successfully.
  Current function value: 50.244967
  Iterations: 83
  Function evaluations: 155
Optimization terminated successfully.
  Current function value: 52.133760
  Iterations: 84
  Function evaluations: 158
Optimization terminated successfully.
  Current function value: 54.230665
  Iterations: 82
  Function evaluations: 156
Optimization terminated successfully.
  Current function value: 56.810668
  Iterations: 90
  Function evaluations: 168
Optimization terminated successfully.
```

```
Current function value: 58.207053
Iterations: 89
Function evaluations: 168
Optimization terminated successfully.
Current function value: 61.041140
Iterations: 87
Function evaluations: 166
Optimization terminated successfully.
Current function value: 64.232456
Iterations: 84
Function evaluations: 159
Optimization terminated successfully.
Current function value: 67.123962
Iterations: 81
Function evaluations: 155
Optimization terminated successfully.
Current function value: 69.299918
Iterations: 85
Function evaluations: 163
Optimization terminated successfully.
Current function value: 70.903428
Iterations: 81
Function evaluations: 152
Optimization terminated successfully.
Current function value: 81.448277
Iterations: 89
Function evaluations: 163
Optimization terminated successfully.
Current function value: 83.804360
Iterations: 83
Function evaluations: 161
Optimization terminated successfully.
Current function value: 86.149082
Iterations: 82
Function evaluations: 160
Optimization terminated successfully.
Current function value: 88.074644
Iterations: 87
Function evaluations: 166
Optimization terminated successfully.
Current function value: 89.930853
Iterations: 79
Function evaluations: 151
Optimization terminated successfully.
Current function value: 92.681328
Iterations: 89
Function evaluations: 167
Optimization terminated successfully.
Current function value: 99.286250
Iterations: 92
Function evaluations: 173
Optimization terminated successfully.
Current function value: 102.287242
Iterations: 84
Function evaluations: 160
Optimization terminated successfully.
Current function value: 103.799865
Iterations: 80
Function evaluations: 154
Optimization terminated successfully.
Current function value: 106.080888
Iterations: 80
Function evaluations: 156
Optimization terminated successfully.
Current function value: 107.605565
Iterations: 79
Function evaluations: 153
Optimization terminated successfully.
Current function value: 113.058277
Iterations: 93
Function evaluations: 178
Optimization terminated successfully.
Current function value: 114.505336
Iterations: 82
Function evaluations: 154
Optimization terminated successfully.
Current function value: 115.913102
Iterations: 96
Function evaluations: 182
Optimization terminated successfully.
Current function value: 117.751392
Iterations: 79
Function evaluations: 151
Optimization terminated successfully.
Current function value: 119.367589
Iterations: 77
Function evaluations: 149
Optimization terminated successfully.
Current function value: 120.875940
```



```
Iterations: 88
Function evaluations: 165
Optimization terminated successfully.
Optimization terminated successfully.
Current function value: 124.737112
Iterations: 85
Function evaluations: 165
Optimization terminated successfully.
Current function value: 129.670497
Iterations: 79
Function evaluations: 150
Optimization terminated successfully.
Current function value: 137.022564
Iterations: 84
Function evaluations: 159
Optimization terminated successfully.
Current function value: 138.590052
Iterations: 85
Function evaluations: 165
Optimization terminated successfully.
Current function value: 140.275050
Iterations: 77
Function evaluations: 146
Optimization terminated successfully.
Current function value: 145.009137
Iterations: 77
Function evaluations: 147
Optimization terminated successfully.
Current function value: 147.811478
Iterations: 84
Function evaluations: 161
Optimization terminated successfully.
Current function value: 150.321737
Iterations: 78
Function evaluations: 150
Optimization terminated successfully.
Current function value: 151.758095
Iterations: 81
Function evaluations: 156
Optimization terminated successfully.
Current function value: 156.753771
Iterations: 75
Function evaluations: 146
Optimization terminated successfully.
Current function value: 158.454630
Iterations: 79
Function evaluations: 155
Optimization terminated successfully.
Current function value: 161.544737
Iterations: 81
Function evaluations: 155
Optimization terminated successfully.
Current function value: 163.644613
Iterations: 80
Function evaluations: 156
Optimization terminated successfully.
Current function value: 165.661547
Iterations: 80
Function evaluations: 156
Optimization terminated successfully.
Current function value: 170.372727
Iterations: 75
Function evaluations: 144
Optimization terminated successfully.
Current function value: 173.232661
Iterations: 83
Function evaluations: 157
Optimization terminated successfully.
Current function value: 174.725255
Iterations: 78
Function evaluations: 150
Optimization terminated successfully.
Current function value: 176.658109
Iterations: 76
Function evaluations: 148
Optimization terminated successfully.
Current function value: 180.110992
Iterations: 82
Function evaluations: 157
Optimization terminated successfully.
Current function value: 183.914366
Iterations: 78
Function evaluations: 149
Optimization terminated successfully.
Current function value: 185.635596
Iterations: 76
Function evaluations: 148
Optimization terminated successfully.
Current function value: 193.930732
```

```
Iterations: 75
Function evaluations: 147
Optimization terminated successfully.
Current function value: 195.423769
Iterations: 84
Function evaluations: 159
Optimization terminated successfully.
Current function value: 200.483857
Iterations: 82
Function evaluations: 156
Optimization terminated successfully.
Current function value: 202.094245
Iterations: 86
Function evaluations: 159
Optimization terminated successfully.
Current function value: 203.595402
Iterations: 83
Function evaluations: 159
Optimization terminated successfully.
Current function value: 205.049057
Iterations: 82
Function evaluations: 155
Optimization terminated successfully.
Current function value: 207.860953
Iterations: 76
Function evaluations: 144
Optimization terminated successfully.
Current function value: 211.460461
Iterations: 81
Function evaluations: 151
Optimization terminated successfully.
Current function value: 214.029520
Iterations: 81
Function evaluations: 152
Optimization terminated successfully.
Current function value: 216.645916
Iterations: 77
Function evaluations: 149
Optimization terminated successfully.
Current function value: 223.085923
Iterations: 77
Function evaluations: 148
Optimization terminated successfully.
Current function value: 224.701822
Iterations: 75
Function evaluations: 146
Optimization terminated successfully.
Current function value: 226.255690
Iterations: 81
Function evaluations: 158
Optimization terminated successfully.
Current function value: 228.090479
Iterations: 75
Function evaluations: 146
Optimization terminated successfully.
Current function value: 230.369384
Iterations: 85
Function evaluations: 158
Optimization terminated successfully.
Current function value: 231.806920
Iterations: 78
Function evaluations: 150
Optimization terminated successfully.
Current function value: 234.619301
Iterations: 74
Function evaluations: 143
Optimization terminated successfully.
Current function value: 236.252667
Iterations: 74
Function evaluations: 143
Optimization terminated successfully.
Current function value: 237.772496
Iterations: 80
Function evaluations: 155
Optimization terminated successfully.
Current function value: 239.912695
Iterations: 78
Function evaluations: 149
Optimization terminated successfully.
Current function value: 241.298176
Iterations: 77
Function evaluations: 150
Optimization terminated successfully.
Current function value: 243.055085
Iterations: 78
Function evaluations: 152
Optimization terminated successfully.
Current function value: 244.628091
Iterations: 84
```

```
Function evaluations: 161
Optimization terminated successfully.
Current function value: 246.030436
Iterations: 85
Function evaluations: 163
Optimization terminated successfully.
Current function value: 247.410376
Current function value: 247.410376
Iterations: 86
Function evaluations: 163
Optimization terminated successfully.
Current function value: 249.646775
Iterations: 80
Function evaluations: 149
Optimization terminated successfully.
Current function value: 253.717529
Iterations: 74
Function evaluations: 143
Optimization terminated successfully.
Current function value: 255.068693
Iterations: 75
Function evaluations: 145
Optimization terminated successfully.
Current function value: 257.714717
Iterations: 80
Function evaluations: 154
Optimization terminated successfully.
Current function value: 259.017250
Iterations: 80
Function evaluations: 155
Optimization terminated successfully.
Current function value: 260.625274
Iterations: 87
Function evaluations: 162
Optimization terminated successfully.
Current function value: 263.544657
Iterations: 79
Function evaluations: 152
Optimization terminated successfully.
Current function value: 268.885058
Iterations: 80
Function evaluations: 150
Optimization terminated successfully.
Current function value: 271.699869
Iterations: 79
Function evaluations: 152
Optimization terminated successfully.
Current function value: 273.029173
Iterations: 78
Function evaluations: 152
Optimization terminated successfully.
Current function value: 274.690860
Iterations: 80
Function evaluations: 151
Optimization terminated successfully.
Current function value: 276.182424
Iterations: 80
Function evaluations: 152
Optimization terminated successfully.
Current function value: 282.041240
Iterations: 78
Function evaluations: 151
Optimization terminated successfully.
Current function value: 283.331904
Iterations: 81
Function evaluations: 154
Optimization terminated successfully.
Current function value: 288.222701
Iterations: 80
Function evaluations: 152
Optimization terminated successfully.
Current function value: 290.788539
Iterations: 86
Function evaluations: 158
Optimization terminated successfully.
Current function value: 295.751863
Iterations: 85
Function evaluations: 162
Optimization terminated successfully.
Current function value: 302.032656
Iterations: 76
Function evaluations: 149
Optimization terminated successfully.
Current function value: 303.756361
Iterations: 72
Function evaluations: 140
Optimization terminated successfully.
Current function value: 306.104533
Iterations: 72
```

```
Function evaluations: 142
Optimization terminated successfully.
Current function value: 309.645428
Iterations: 78
Function evaluations: 150
Optimization terminated successfully.
Current function value: 310.998849
Iterations: 76
Function evaluations: 148
Optimization terminated successfully.
Current function value: 312.741050
Iterations: 82
Function evaluations: 152
Optimization terminated successfully.
Current function value: 316.785037
Iterations: 72
Function evaluations: 140
Optimization terminated successfully.
Current function value: 318.246655
Iterations: 72
Function evaluations: 141
Optimization terminated successfully.
Current function value: 319.990919
Iterations: 73
Function evaluations: 142
Optimization terminated successfully.
Current function value: 324.825810
Iterations: 84
Function evaluations: 155
Optimization terminated successfully.
Current function value: 333.450131
Iterations: 75
Function evaluations: 143
Optimization terminated successfully.
Current function value: 335.966797
Iterations: 79
Function evaluations: 151
Optimization terminated successfully.
Current function value: 337.402793
Iterations: 76
Function evaluations: 146
Optimization terminated successfully.
Current function value: 340.020481
Iterations: 90
Function evaluations: 163
Optimization terminated successfully.
Current function value: 341.502172
Iterations: 89
Function evaluations: 164
Optimization terminated successfully.
Current function value: 343.701537
Iterations: 86
Function evaluations: 157
Optimization terminated successfully.
Current function value: 347.106162
Iterations: 88
Function evaluations: 162
Optimization terminated successfully.
Current function value: 349.540678
Iterations: 81
Function evaluations: 152
Optimization terminated successfully.
Current function value: 354.252588
Iterations: 84
Function evaluations: 161
Optimization terminated successfully.
Current function value: 355.900194
Iterations: 103
Function evaluations: 192
Optimization terminated successfully.
Current function value: 358.386896
Iterations: 88
Function evaluations: 170
Optimization terminated successfully.
Current function value: 363.858308
Iterations: 79
Function evaluations: 150
Optimization terminated successfully.
Current function value: 365.286566
Iterations: 77
Function evaluations: 149
Optimization terminated successfully.
Current function value: 366.900914
Iterations: 74
Function evaluations: 143
Optimization terminated successfully.
Current function value: 368.898213
Iterations: 89
```

```
Function evaluations: 161
Optimization terminated successfully.
Current function value: 370.262323
Iterations: 85
Function evaluations: 156
Optimization terminated successfully.
Current function value: 379.016732
Iterations: 79
Function evaluations: 151
Optimization terminated successfully.
Current function value: 380.842308
Iterations: 81
Iterations: 81
Function evaluations: 155
Optimization terminated successfully.
Current function value: 383.228490
Iterations: 84
Function evaluations: 157
Optimization terminated successfully.
Current function value: 384.785304
Iterations: 80
Function evaluations: 152
Optimization terminated successfully.
Current function value: 386.200290
Iterations: 78
Function evaluations: 151
Optimization terminated successfully.
Current function value: 388.861944
Iterations: 79
Function evaluations: 151
Optimization terminated successfully.
Current function value: 390.290678
Iterations: 76
Function evaluations: 145
Optimization terminated successfully.
Current function value: 393.179383
Iterations: 85
Function evaluations: 163
Optimization terminated successfully.
Current function value: 394.721305
Iterations: 84
Function evaluations: 159
Optimization terminated successfully.
Current function value: 396.150884
Iterations: 82
Function evaluations: 157
Optimization terminated successfully.
Current function value: 402.837239
Iterations: 75
Function evaluations: 145
Optimization terminated successfully.
Current function value: 404.373868
Iterations: 81
Function evaluations: 156
Optimization terminated successfully.
Current function value: 407.602372
Iterations: 85
Function evaluations: 156
Optimization terminated successfully.
Current function value: 410.008722
Iterations: 84
Function evaluations: 157
Optimization terminated successfully.
Current function value: 411.399714
Iterations: 74
Function evaluations: 143
Optimization terminated successfully.
Current function value: 413.591007
Iterations: 84
Function evaluations: 154
Optimization terminated successfully.
Current function value: 415.467720
Iterations: 86
Function evaluations: 158
Optimization terminated successfully.
Current function value: 416.925867
Iterations: 84
Function evaluations: 155
Optimization terminated successfully.
Current function value: 419.522106
Iterations: 91
Function evaluations: 168
Optimization terminated successfully.
Current function value: 421.936961
Iterations: 84
Function evaluations: 155
Optimization terminated successfully.
Current function value: 424.538795
Iterations: 75
```

```
Function evaluations: 143
Optimization terminated successfully.
Current function value: 426.327458
Iterations: 80
Function evaluations: 154
Optimization terminated successfully.
Current function value: 430.445037
Iterations: 77
Function evaluations: 151
Optimization terminated successfully.
Current function value: 440.428913
Iterations: 73
Function evaluations: 143
Optimization terminated successfully.
Current function value: 441.784183
Iterations: 76
Function evaluations: 146
Optimization terminated successfully.
Current function value: 447.248164
Iterations: 85
Function evaluations: 161
Optimization terminated successfully.
Current function value: 448.826613
Iterations: 81
Function evaluations: 153
Optimization terminated successfully.
Current function value: 450.261298
Iterations: 85
Function evaluations: 161
Optimization terminated successfully.
Current function value: 451.674732
Iterations: 84
Function evaluations: 161
Optimization terminated successfully.
Current function value: 453.009800
Iterations: 83
Function evaluations: 157
Optimization terminated successfully.
Current function value: 456.343129
Iterations: 87
Function evaluations: 166
Optimization terminated successfully.
Current function value: 459.355816
Iterations: 75
Function evaluations: 145
Optimization terminated successfully.
Current function value: 462.712419
Iterations: 86
Function evaluations: 162
Optimization terminated successfully.
Current function value: 464.638954
Iterations: 83
Function evaluations: 159
Optimization terminated successfully.
Current function value: 466.793784
Iterations: 86
Function evaluations: 162
Optimization terminated successfully.
Current function value: 472.890358
Iterations: 77
Function evaluations: 149
Optimization terminated successfully.
Current function value: 477.349859
Iterations: 91
Function evaluations: 168
Optimization terminated successfully.
Current function value: 483.603146
Iterations: 77
Function evaluations: 148
Optimization terminated successfully.
Current function value: 485.003674
Iterations: 79
Function evaluations: 149
Optimization terminated successfully.
Current function value: 486.664846
Iterations: 78
Function evaluations: 149
Optimization terminated successfully.
Current function value: 496.317641
Iterations: 77
Function evaluations: 151
Optimization terminated successfully.
Current function value: 498.521531
Iterations: 79
Function evaluations: 153
Optimization terminated successfully.
Current function value: 501.882182
Iterations: 77
```

```
Function evaluations: 151
Optimization terminated successfully.
Current function value: 503.383302
Iterations: 77
Function evaluations: 151
Optimization terminated successfully.
Current function value: 504.991433
Iterations: 78
Function evaluations: 152
Optimization terminated successfully.
Current function value: 509.052355
Iterations: 81
Function evaluations: 153
Optimization terminated successfully.
Current function value: 511.747097
Iterations: 81
Function evaluations: 153
Optimization terminated successfully.
Current function value: 513.830133
Iterations: 82
Function evaluations: 155
Optimization terminated successfully.
Current function value: 515.676904
Iterations: 79
Function evaluations: 150
Optimization terminated successfully.
Current function value: 517.201032
Iterations: 86
Function evaluations: 163
Optimization terminated successfully.
Current function value: 521.005817
Iterations: 85
Function evaluations: 160
Optimization terminated successfully.
Current function value: 522.718773
Iterations: 89
Function evaluations: 172
Optimization terminated successfully.
Current function value: 524.126741
Iterations: 90
Function evaluations: 172
Optimization terminated successfully.
Current function value: 526.233875
Iterations: 84
Function evaluations: 158
Optimization terminated successfully.
Current function value: 527.867788
Iterations: 88
Function evaluations: 161
Optimization terminated successfully.
Current function value: 533.017629
Iterations: 77
Function evaluations: 148
Optimization terminated successfully.
Current function value: 535.002388
Iterations: 76
Function evaluations: 147
Optimization terminated successfully.
Current function value: 536.971966
Iterations: 77
Function evaluations: 150
Optimization terminated successfully.
Current function value: 538.406220
Iterations: 78
Function evaluations: 151
Optimization terminated successfully.
Current function value: 539.799025
Iterations: 77
Function evaluations: 150
Optimization terminated successfully.
Current function value: 543.127376
Iterations: 78
Function evaluations: 150
Optimization terminated successfully.
Current function value: 546.051274
Iterations: 78
Function evaluations: 151
Optimization terminated successfully.
Current function value: 548.515494
Iterations: 79
Function evaluations: 153
Optimization terminated successfully.
Current function value: 552.751043
Iterations: 79
Function evaluations: 151
Optimization terminated successfully.
Current function value: 554.131287
Iterations: 79
```

```

Function evaluations: 149
Optimization terminated successfully.
Current function value: 557.188955
Iterations: 85
Function evaluations: 157

```

-----FileNotFoundError

```

Traceback (most recent call last)
<ipython-input-45-6a1108c4c769> in <cell line: 32>()

```

```

import pandas as pd

```

```

import numpy as np

```

```

a=pd.read_csv("train.csv",sep=';')

```

```

print(a)

```

```

p=a[['hour']]

```

```

m=np.mean(p)

```

```

sd=np.std(p)#sigma value

```

```

var=np.var(p)#sigma square value

```

```

m, sd, var

```

```

id year hour season holiday workingday weather temp atemp \
0 3 2012 23 3 0 0 2 23.78 27.275
1 4 2011 8 3 0 0 1 27.88 31.820
2 5 2012 2 1 0 1 1 20.50 24.240
3 7 2011 20 3 0 1 3 25.42 28.790
4 8 2011 17 3 0 1 3 26.24 28.790
...
7684 10882 2012 18 1 0 1 1 13.94 15.150
7685 10883 2012 3 1 0 1 1 9.02 11.365
7686 10884 2012 15 2 0 0 1 21.32 25.000
7687 10885 2011 19 4 0 1 1 12.30 14.395
7688 10886 2012 21 3 0 1 1 30.34 34.850

```

```

humidity windspeed count
0 73 11.0014 133
1 57 0.0000 132
2 59 0.0000 19
3 83 19.9995 58
4 89 0.0000 285
...
7684 42 22.0028 457
7685 51 11.0014 1
7686 19 27.9993 626
7687 45 15.0013 217
7688 66 7.0015 381

```

```

[7689 rows x 12 columns]
(11.56535310183379, 6.915326938018648, 47.82174665968637)

```

```

t=np.array(a['temp'])
tm=np.mean(t) tsd=np.std(t)#sigma
value/std.dev
tvar=np.var(t)#variance x=29
l=np.log(np.sqrt(2*3.14))
e=np.log(tsd) #std.dev
f=(x-tm)**2 #mean
g=2*(tvar**2) #variance
h=f/g i=-l-e print(i-h)
-2.9860025235468406

```

```

import pandas as pd import numpy
as np
a=pd.read_csv("test.csv",sep=';')

```



```
print(a) p=a['hour'] m=np.mean(p)
sd=np.std(p)#sigma value
var=np.var(p)#sigma square value
m, sd, var
```

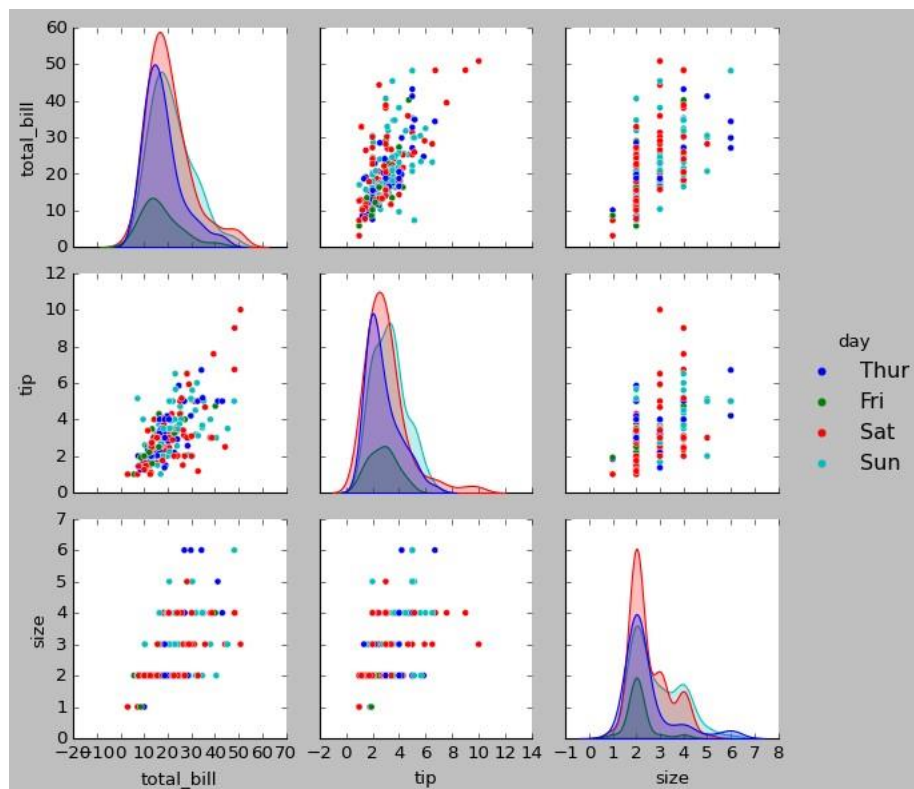
```

Unnamed: 0  year  hour  season  holiday  workingday  weather  temp  \
0          1   2012   21       3         0           0       1  29.52
1          2   2012    3       2         0           0       1  23.78
2          6   2011   10       1         0           1       3  16.40
3         14   2012   19       1         0           1       1  13.94
4         17   2011   23       3         0           1       2  26.24  ...
...
3191      10868  2012    9       4         0           1       1  16.40
3192      10871  2011   12       4         0           1       3  18.04
3193      10872  2011   19       3         0           1       1  30.34
3194      10873  2012    0       4         0           1       2  13.12  3195      10874  2012   10       1       0       1
3  12.30

      atemp  humidity  windspeed
0    34.850        79     6.0032
1    27.275        83     0.0000
2    20.455         0    11.0014
3    15.150        46    19.9995
4    30.305        73    11.0014  ...
3191  20.455        87     6.0032
3192  21.970       100     8.9981
3193  34.090        55    16.9979
3194  16.665        66     7.0015
3195  14.395        87    16.9979
```

```
[3196 rows x 11 columns] (11.482478097622028,
6.915772969192248, 47.82791576141015)
```

```
import seaborn import
matplotlib.pyplot as plt df =
seaborn.load_dataset('tips')
seaborn.pairplot(df, hue='day')
plt.show()
```



Linear Regrssion on US Housing Price

Linear Regrssion on US Housing Price

In statistics, linear regression is a linear approach for modeling the relationship between a scalar dependent variable y and one or more explanatory variables (or independent variables) denoted X . The case of one explanatory variable is called simple linear regression. For more than one explanatory variable, the process is called multiple linear regression.

Linear regression models are often fitted using the least squares approach, but they may also be fitted in other ways, such as by minimizing the "lack of fit" in some other norm (as with least absolute deviations regression), or by minimizing a penalized version of the least squares loss function as in ridge regression (L_2 -norm penalty) and lasso (L_1 -norm penalty). Conversely, the least squares approach can be used to fit models that are not linear models. Thus, although the terms "least squares" and "linear model" are closely linked, they are not synonymous.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

df = pd.read_csv("USA_Housing.csv")
df.head()
```

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price	Address
0	79545.458574	5.682861	7.009188	4.09	23086.800503	1.059034e+06	208 Michael F 674\nLaura

Check basic info on the data set

'info()' method to check the data types and number

```
df.info(verbose=True)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 7 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Avg. Area Income                      5000 non-null   float64
1   Avg. Area House Age                   5000 non-null   float64
2   Avg. Area Number of Rooms             5000 non-null   float64
3   Avg. Area Number of Bedrooms          5000 non-null   float64
4   Area Population                       5000 non-null   float64
5   Price                                5000 non-null   float64
6   Address                               5000 non-null   object
dtypes: float64(6), object(1)
memory usage: 273.6+ KB
```

'describe()' method to get the statistical summary of the various features of the data set

```
df.describe(percentiles=[0.1,0.25,0.5,0.75,0.9])
```

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price
count	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000	5.000000e+03
mean	68583.108984	5.977222	6.987792	3.98133036163	516039.516039	1.232073e+06

28/10/2023, 18:18

StatML-Lab03.ipynb - Colaboratory

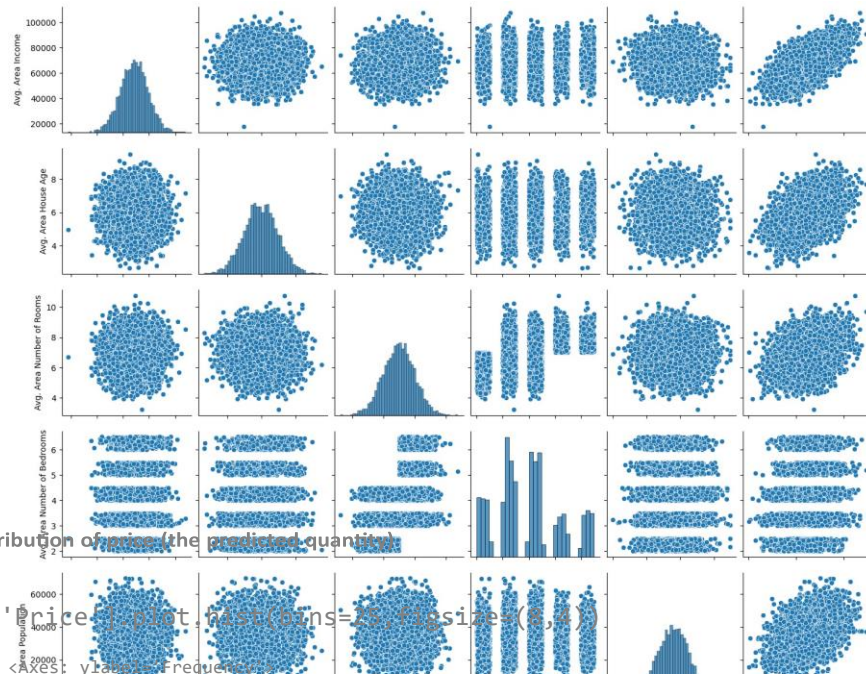
std	10657.991214	0.991456	1.005833	1.234137	9925.650114	3.531176e+05			
'columns' method to get the names of the columns (features)min						172.610686	1.593866e+04		
	17796.631190	2.644304	3.236194	2.000000	23502.845262	7.720318e+05			
10%	55047.633980	4.697755	5.681951	2.310000	29403.928702	9.975771e+05			
df.columns						25%	61480.562388	5.322283	6.299250
Index(['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms', 'Avg. Area Number of Bedrooms', 'Area Population', 'Price', 'Address'],						50%	68804.286404	5.970429	7.002902
dtype=object)						75%	75783.338666	6.650808	7.665871
						90%	82084.488282	7.242078	8.274222
						max	106579.991214	10.005833	12.341370

Basic plotting and visualization on the data set

Pairplots using seaborn


```
sns.pairplot(df)
```

```
<seaborn.axisgrid.PairGrid at 0x7e90c862b340>
```




Distribution of price (the predicted quantity)

```
df['price'].plot.hist(bins=25, figsize=(8,4))
```



```
df['Price'].plot.density()
```

```
<Axes: ylabel='Density'>
```



Correlation matrix and heatmap

```
df.corr()
```

```
<ipython-input-9-2f6f6606aa2c>:1: FutureWarning: The default value of numeric_only in
df.corr()
```

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price
Avg. Area Income	1.000000	-0.002007	-0.011032	0.019788	-0.016234	0.639734
Avg. Area House Age	-0.002007	1.000000	-0.009428	0.006149	-0.018743	0.452543
Avg. Area Number of Rooms	-0.011032	-0.009428	1.000000	0.462695	0.002040	0.335664
Avg. Area Number of Bedrooms				1.000000	0.002040	0.335664
Area Population				0.002040	1.000000	0.335664
Price				0.335664	0.335664	1.000000

```
plt.figure(figsize=(10,7))
sns.heatmap(df.corr(),annot=True,linewidths=2)
```

```
hon-input-10-73d88c5a3f1a>:2: FutureWarning: The default value of numeric_only in
sns.heatmap(df.corr(),annot=True,linewidths=2)
: >
```



Feature and variable sets

Make a list of data frame column names

```
l_column = list(df.columns) # Making a list out of column names
len_feature = len(l_column) # Length of column vector list
l_column
['Avg. Area Income',
 'Avg. Area House Age',
 'Avg. Area Number of Rooms',
```

```
'Avg. Area Number of Bedrooms',
'Area Population',
'Price',
'Address']
```

Put all the numerical features in X and Price in y, ignore Address which is string for linear regression

```
X = df[l_column[0:len_feature-2]]
y = df[l_column[len_feature-2]]
```

```
print("Feature set size:",X.shape)
print("Variable set size:",y.shape)
```

```
Feature set size: (5000, 5)
Variable set size: (5000,)
```

```
X.head()
```

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population
0	79545.458574	5.682861	7.009188	4.09	23086.800503
1	79248.642455	6.002900	6.730821	3.09	40173.072174
2	61287.067179	5.865890	8.512727	5.13	36882.159400
3	63345.240046	7.188236	5.586729	3.26	34310.242831

```
y.head()
```

```
0    1.059034e+06
1    1.505891e+06
2    1.058988e+06
3    1.260617e+06
4    6.309435e+05
Name: Price, dtype: float64
```

Test-train split

Import `train_test_split` function from `scikit-learn`

```
from sklearn.model_selection import train_test_split
```

Create X and y train and test splits in one command using a split ratio and a random seed

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=123)
```

Check the size and shape of train/test splits (it should be in the ratio as per `test_size` parameter above)

```
print("Training feature set size:",X_train.shape)
print("Test feature set size:",X_test.shape)
print("Training variable set size:",y_train.shape)
print("Test variable set size:",y_test.shape)
```

```
Training feature set size: (3500, 5)
Test feature set size: (1500, 5)
Training variable set size: (3500,)
Test variable set size: (1500,)
```

Model fit and training

Import linear regression model estimator from `scikit-learn` and instantiate

```
from sklearn.linear_model import LinearRegression
from sklearn import metrics
lm = LinearRegression() # Creating a Linear Regression object 'lm'
```

Fit the model on to the instantiated object itself

```
lm.fit(X_train,y_train) # Fit the linear model on to the 'lm' object itself i.e. no need to set
```

```
LinearRegression()
LinearRegression()
```

Check the intercept and coefficients and put them in a DataFrame

```
print("The intercept term of the linear model:", lm.intercept_)
```

```
The intercept term of the linear model: -2631028.9017454907
```

```
print("The coefficients of the linear model:", lm.coef_)
```

```
The coefficients of the linear model: [2.15976020e+01 1.65201105e+05 1.19061464e+05 3.21258561e+03
1.52281212e+01]
```

```
#idict = {'Coefficients':lm.intercept_}
#idf = pd.DataFrame(data=idict,index=['Intercept'])
cdf = pd.DataFrame(data=lm.coef_, index=X_train.columns, columns=["Coefficients"])
#cdf=pd.concat([idf,cdf], axis=0)
cdf
```

	Coefficients
Avg. Area Income	21.597602
Avg. Area House Age	165201.104954
Avg. Area Number of Rooms	119061.463868
Avg. Area Number of Bedrooms	3212.585606
Area Population	15.228121

Calculation of standard errors and t-statistic for the coefficients

```
n=X_train.shape[0]
k=X_train.shape[1]
dfN = n-k
train_pred=lm.predict(X_train)
train_error = np.square(train_pred - y_train)
sum_error=np.sum(train_error) se=[0,0,0,0,0]
for i in range(k):
    r = (sum_error/dfN)
    r = r/np.sum(np.square(X_train[list(X_train.columns)[i]]-X_train[list(X_train.columns)[i]]).m
se[i]=np.sqrt(r) cdf['Standard Error']=se
cdf['t-statistic']=cdf['Coefficients']/cdf['Standard Error']
cdf
```

	Coefficients	Standard Error	t-statistic
Avg. Area Income	21.597602	0.160361	134.681505
Avg. Area House Age	165201.104954	1722.412068	95.912649
Avg. Area Number of Rooms	119061.463868	1696.546476	70.178722
Avg. Area Number of Bedrooms	3212.585606	1376.451759	2.333962
Area Population	15.228121	0.169882	89.639472

```
print("Therefore, features arranged in the order of importance for predicting the house price\n"
l=list(cdf.sort_values('t-statistic',ascending=False).index) print(' > \n'.join(l))
```

```
Therefore, features arranged in the order of importance for predicting the house price -----
-----
Avg. Area Income >
Avg. Area House Age >
Area Population >
Avg. Area Number of Rooms >
Avg. Area Number of Bedrooms
```

```

l=list(cdf.index)
from matplotlib import gridspec
fig = plt.figure(figsize=(18, 10))
gs = gridspec.GridSpec(2,3)
#f, ax = plt.subplots(nrows=1,ncols=len(l), sharey=True)
ax0 = plt.subplot(gs[0])
ax0.scatter(df[l[0]],df['Price'])
ax0.set_title(l[0]+" vs. Price", fontdict={'fontsize':20})

ax1 = plt.subplot(gs[1])
ax1.scatter(df[l[1]],df['Price'])
ax1.set_title(l[1]+" vs. Price",fontdict={'fontsize':20})

ax2 = plt.subplot(gs[2])
ax2.scatter(df[l[2]],df['Price'])
ax2.set_title(l[2]+" vs. Price",fontdict={'fontsize':20})

ax3 = plt.subplot(gs[3])
ax3.scatter(df[l[3]],df['Price'])
ax3.set_title(l[3]+" vs. Price",fontdict={'fontsize':20})

```

```

ax4 = plt.subplot(gs[4])
ax4.scatter(df[l[4]],df['Price'])
ax4.set_title(l[4]+" vs. Price",fontdict={'fontsize':20})
Text(0.5, 1.0, 'Area Population vs. Price')

```

```

print("R-squared value of this fit:",round(metrics.r2_score(y_train,train_pred),3))

```

R-squared value of this fit: 0.917
 Prediction, error estimate, and regression evaluation matrices
 Prediction using the lm model

```

predictions = lm.predict(X_test)
print ("Type of the predicted object:", type(predictions))
print ("Size of the predicted object:", predictions.shape)

```

```

Type of the predicted object: <class 'numpy.ndarray'>
Size of the predicted object: (1500,)

```

Scatter plot of predicted price and y_test set to see if the data fall on a 45 degree straight line

```


plt.figure(figsize=(10,7))
plt.title("Actual vs. predicted house prices",fontsize=25)
plt.xlabel("Actual test set house prices",fontsize=18)
plt.ylabel("Predicted house prices", fontsize=18)
plt.scatter(x=y_test,y=predictions)

```

```

<matplotlib.collections.PathCollection at 0x7e90bf811d20>

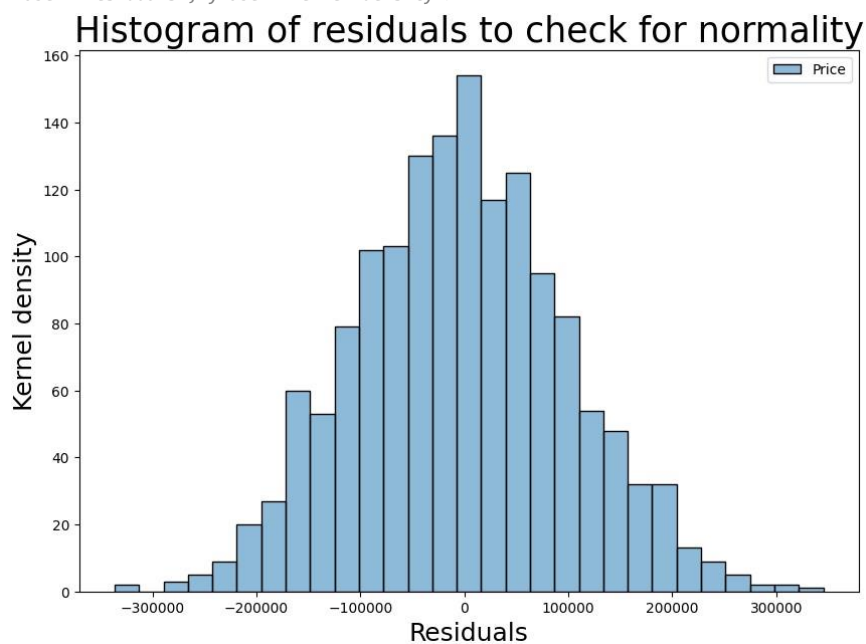
```

Plotting histogram of the residuals i.e. predicted errors (expect a normally distributed pattern)

```
plt.figure(figsize=(10,7))
plt.title("Histogram of residuals to check for normality",fontsize=25)
plt.xlabel("Residuals",fontsize=18)      plt.ylabel("Kernel density",      fontsize=18)
sns.histplot([y_test-predictions])
```

```
<Axes: title={'center': 'Histogram of residuals to check for normality'},
      xlabel='Residuals', ylabel='Kernel density'>
```

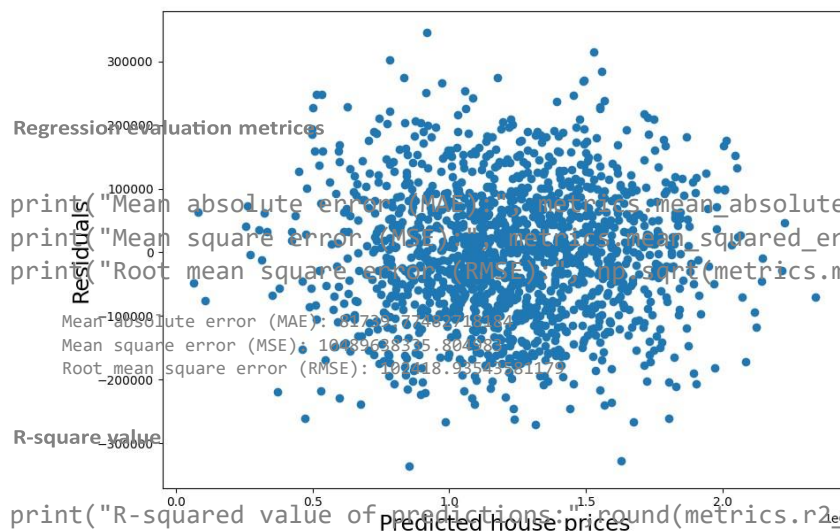


Scatter plot of residuals and predicted values (Homoscedasticity)

```
plt.figure(figsize=(10,7))
plt.title("Residuals vs. predicted values plot (Homoscedasticity)\n",fontsize=25)
plt.xlabel("Predicted house prices",fontsize=18) plt.ylabel("Residuals",
fontsize=18) plt.scatter(x=predictions,y=y_test-predictions)
```

<matplotlib.collections.PathCollection at 0x7e90bf7309d0>

Residuals vs. predicted values plot (Homoscedasticity)



```
print("Mean absolute error (MAE):", metrics.mean_absolute_error(y_test,predictions))
print("Mean square error (MSE):", metrics.mean_squared_error(y_test,predictions))
print("Root mean square error (RMSE):", np.sqrt(metrics.mean_squared_error(y_test,predictions)))
```

Mean absolute error (MAE): 10795.746310434
Mean square error (MSE): 103963635.804085
Root mean square error (RMSE): 10196.71179

```
print("R-squared value of predictions: round(metrics.r2_score(y_test,predictions),3))
```

R-squared value of predictions: 0.919

```
#compute minmax value for observed price and expected price
```

```
import numpy as np min=np.min(predictions/6000)
```

```
max=np.max(predictions/12000) print(min, max)
```

10.57339854753646 195.14363973516853


```
#Compute MinMax value for Price=100
```

```
L = (100 - min)/(max - min)
```

```
L plt.hist(L)
```

```
(array([0., 0., 0., 0., 0., 1., 0., 0., 0., 0.]),
 array([-0.01548743,  0.08451257,  0.18451257,  0.28451257,  0.38451257,
         0.48451257,  0.58451257,  0.68451257,  0.78451257,  0.88451257,
         0.98451257])),
```

<BarContainer object of 10 artists>)



Logistic Regression with Titanic data set

Import packages and dataset

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import seaborn as sns
import nbconvert

dataframe = pd.read_csv("titanic_train.csv")
dataframe.head()
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	F
0	1	0	3	Braund, Mr. Owen	male	22.0	1	0	A/5 21171	7.2
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2

Check basic info about the data set including missing value

```
dataframe.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  -
0   PassengerId  891 non-null    int64
1   Survived     891 non-null    int64
2   Pclass       891 non-null    int64
3   Name         891 non-null    object
4   Sex          891 non-null    object
5   Age          714 non-null    float64
6   SibSp        891 non-null    int64
7   Parch        891 non-null    int64
8   Ticket       891 non-null    object
9   Fare         891 non-null    float64
10  Cabin        204 non-null    object
11  Embarked     889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

```
d=dataframe.describe()
d
```

	PassengerId	Survived	Pclass	Age	SibSp	Parch	F
count	891.000000	891.000000	891.000000	714.000000	891.000000	891.000000	891.000
mean	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594	32.204
std	257.353842	0.486592	0.836071	14.526497	1.102743	0.806057	49.693
min	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000	0.000
25%	223.500000	0.000000	2.000000	20.125000	0.000000	0.000000	7.910
50%	446.000000	0.000000	3.000000	28.000000	0.000000	0.000000	14.454
75%	668.500000	1.000000	3.000000	38.000000	1.000000	0.000000	31.000
max	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000	512.329

Exploratory analysis and plots

Plot a bar diagram to check the number of numeric entries

From the bar diagram, it shows that there are some age entries missing as the number of count for 'Age' is less than the other counts. We can do some impute/transformation of the data to fill-up the missing entries.


```
result = dataframe.dtypes
result
```

```
PassengerId    int64
Survived        int64
Pclass         int64
Name           object
Sex            object
Age            float64
SibSp          int64
Parch          int64
Ticket         object
Fare           float64
Cabin          object
Embarked       object
dtype: object
```

```
dT=d.T
```



```
dT.plot.bar(y='count') plt.title("Bar plot of the count of numeric
features", fontsize=17)
```

```
Text(0.5, 1.0, 'Bar plot of the count of numeric features')
```



Check the relative size of survived and not-survived

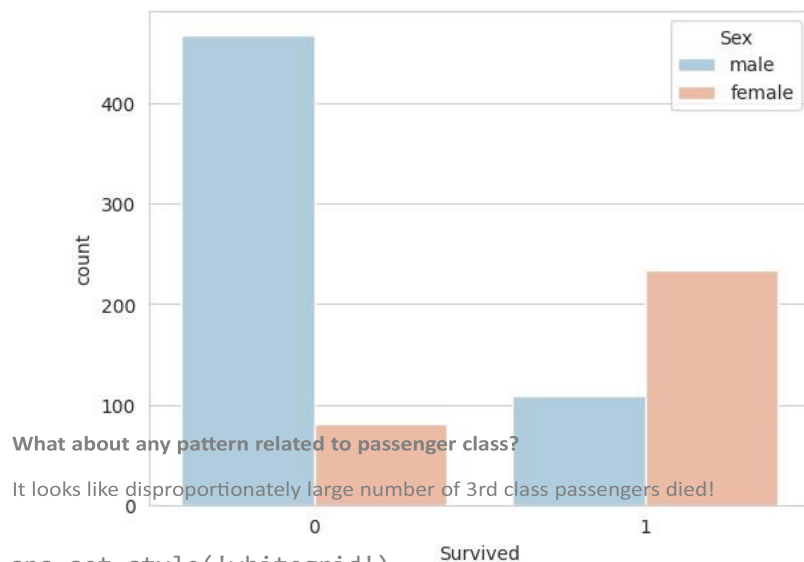
```
sns.set_style('whitegrid')
sns.countplot(x='Survived', data=dataframe, palette='RdBu_r')
sns.pairplot(dataframe)
```




Is there a pattern for the survivability based on sex?

It looks like more female survived than males!

```
sns.set_style('whitegrid')
sns.countplot(x='Survived', hue='Sex', data=dataframe, palette='RdBu_r')
<Axes: xlabel='Survived', ylabel='count'>
```



```
sns.set_style('whitegrid')
sns.countplot(x='Survived',hue='Pclass',data=dataframe,palette='rainbow')
<Axes: xlabel='Survived', ylabel='count'>
```



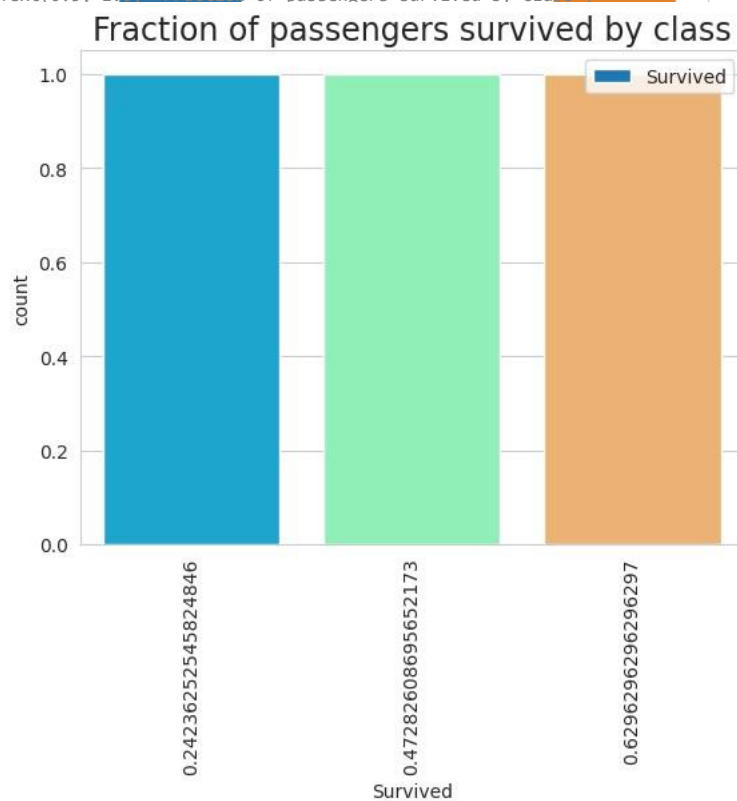
```
sns.set_style('whitegrid')
plt.figure(figsize=(6, 4))
sns.countplot(data=dataframe, x='Sex', hue='Survived')
plt.title("Survival Count by Gender")
plt.xlabel("Gender") plt.ylabel("Count") plt.show()
```

Survival Count by Gender

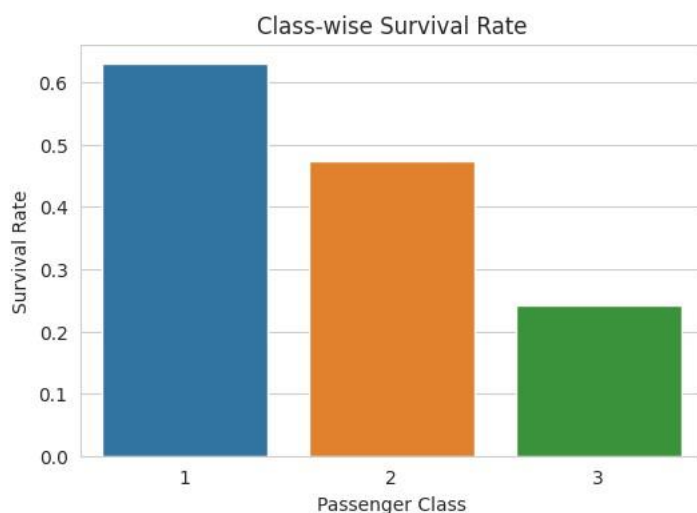
Following code extracts and plots the fraction of passenger count that survived, by each class

```
f_class_survived = dataframe.groupby('Pclass')['Survived'].mean()
f_class_survived = pd.DataFrame(f_class_survived)
f_class_survived
f_class_survived.plot.bar(y='Survived')
sns.countplot(x='Survived', data=f_class_survived, palette='rainbow')
plt.title("Fraction of passengers survived by class", fontsize=17)
```

Text(0.5, 1.0, 'Fraction of passengers survived by class')



```
class_survival = dataframe.groupby('Pclass')['Survived'].mean()
plt.figure(figsize=(6, 4))
sns.barplot(x=class_survival.index, y=class_survival.values)
plt.title("Class-wise Survival Rate") plt.xlabel("Passenger Class") plt.ylabel("Survival Rate") plt.show()
```




What about any pattern related to having sibling and spouse?

It looks like there is a weak trend that chance of survivability increased if there were more number of sibling or spouse

```
sns.set_style('whitegrid')
sns.countplot(x='Survived',hue='SibSp',data=dataframe,palette='rainbow')
```


<Axes: xlabel='Survived', ylabel='count'>



How does the overall age distribution look like?

```
plt.xlabel("Age of the passengers",fontsize=18)
plt.ylabel("Count",fontsize=18)
plt.title("Age histogram of the passengers",fontsize=22)
#train['Age'].hist(bins=30,color='darkred',alpha=0.7,figsize=(10,6))
dataframe['Age'].hist()
```


<Axes: title={'center': 'Age histogram of the passengers'}, xlabel='Age of the passengers', ylabel='Count'>



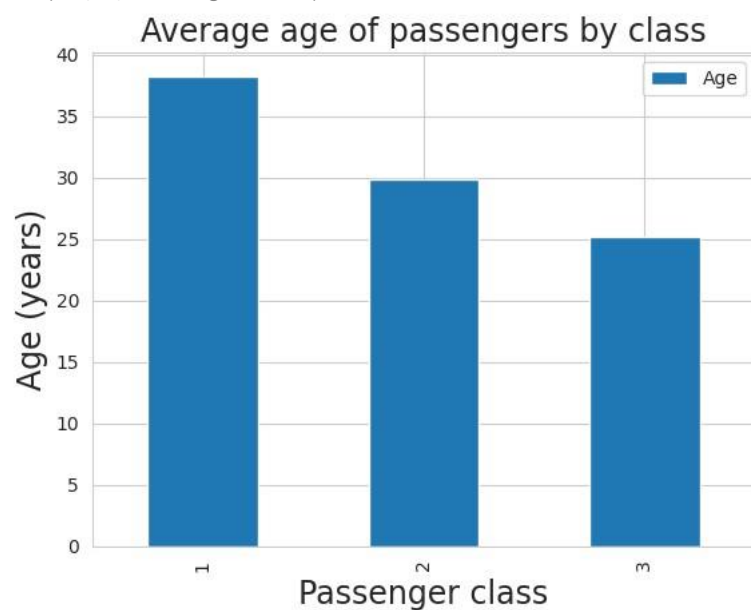
How does the age distribution look like across passenger class?

It looks like that the average age is different for three classes and it generally decreases from 1st class to 3rd class.

```
plt.figure(figsize=(12, 10))
plt.xlabel("Passenger Class",fontsize=18)
plt.ylabel("Age",fontsize=18)
sns.boxplot(x='Pclass',y='Age',data=dataframe,palette='winter')
<Axes: xlabel='Pclass', ylabel='Age'>
```

```
f_class_Age=dataframe.groupby('Pclass')['Age'].mean()
f_class_Age = pd.DataFrame(f_class_Age)
f_class_Age.plot.bar(y='Age')
plt.title("Average age of passengers by class",fontsize=17)
plt.ylabel("Age (years)", fontsize=17)
plt.xlabel("Passenger class", fontsize=17)
Text(0.5, 0, 'Passenger class')
```



Data wrangling (impute and drop)

- Impute age (by averaging)
- Drop unnecessary features
- Convert categorical features to dummy variables

Define a function to impute (fill-up missing values) age feature

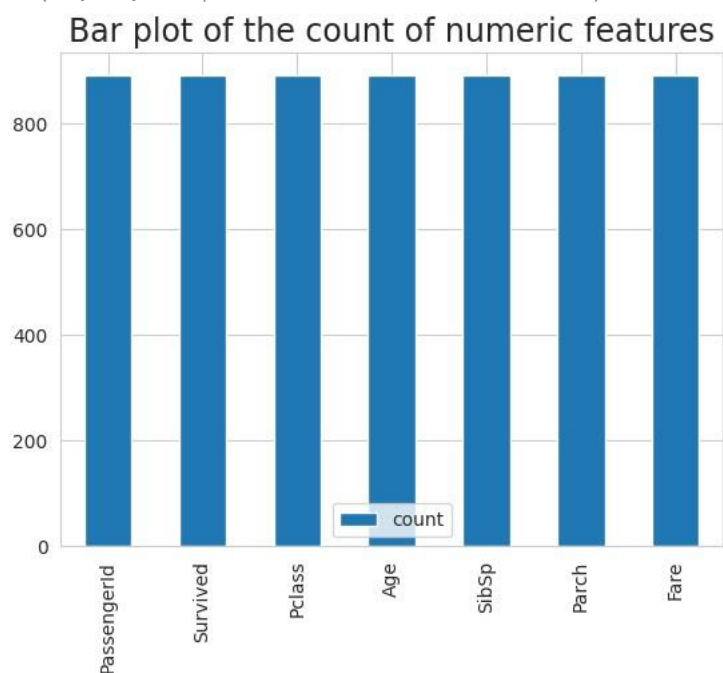
```
a=list(dataframe['Age']);
def impute_age(cols):
    Age = cols[0]    Pclass =
    cols[1]    if
    pd.isnull(Age):    if
    Pclass == 1:        return
    a[0]        elif Pclass ==
    2:            return a[2]
    else:
        return a[2]

    else:
    return Age
```

Apply the above-defined function and plot the count of numeric features

```
dataframe['Age'] = dataframe[['Age', 'Pclass']].apply(impute_age,axis=1)
d=dataframe.describe() dT=d.T dT.plot.bar(y='count') plt.title("Bar
plot of the count of numeric features",fontsize=17)
```

```
Text(0.5, 1.0, 'Bar plot of the count of numeric features')
```



Drop the 'Cabin' feature and any other null value

```
dataframe.drop('Cabin',axis=1,inplace=True)
dataframe.dropna(inplace=True)
dataframe.head()
```

PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	F
0	1	0	3	Braund,					
			Mr. Owen	male	22.0	1	0	A/5 21171	7.2
			Harris						

Cumings,
 Mrs. John
 (Florence
 female 38.0
 PC 17599 71.2

```
dataframe.drop(['PassengerId', 'Name', 'Ticket'], axis=1, inplace=True)
dataframe.head()
```

Convert categorical feature like 'Sex' and 'Embarked' to dummy variables

Use pandas 'get_dummies()' function

```
sex = pd.get_dummies(dataframe['Sex'], drop_first=True)
embark = pd.get_dummies(dataframe['Embarked'], drop_first=True)
```

Now drop the 'Sex' and 'Embarked' columns and concatenate the new dummy variables

```
dataframe.drop(['Sex', 'Embarked'], axis=1, inplace=True)
dataframe = pd.concat([dataframe, sex, embark], axis=1)
dataframe.head()
```

This data set is now ready for logistic regression analysis!

Logistic Regression model t and prediction

Let's start by splitting our data into a training set and test set (there is another test.csv file that you can play around with in case you want to use all this data for training).

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(dataframe.drop('Survived', axis=1),
dataframe['Survived'], test_size=0.30,
random_state=111)
```

F1-score as a function of regularization (penalty) parameter

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report
import matplotlib.pyplot as plt
```

```
nsimu = 201
penalty = [0]
* nsimu logmodel = [0] *
nsimu predictions = [0] *
nsimu class_report = [0] *
nsimu f1 = [0] * nsimu
```

```
for i in range(1, nsimu + 1):
    logmodel[i] = LogisticRegression(C=i / 1000, tol=1e-4, max_iter=int(1e6), n_jobs=4)
    logmodel[i].fit(X_train, y_train)
    predictions[i] = logmodel[i].predict(X_test)
    class_report[i] = classification_report(y_test, predictions[i])
l = class_report[i].split()
f1[i] = l[len(l) - 1]
penalty[i] = 1000 / i
plt.scatter(penalty[1:len(penalty) - 2],
f1[1:len(f1) - 2])
plt.title("F1-score vs. regularization
parameter", fontsize=20)
plt.xlabel("Penalty parameter",
fontsize=17)
plt.ylabel("F1-score on test data", fontsize=17)
plt.show()
```

F1-score as a function of test set size (fraction)

```

nsimu=101
class_report = [0]*nsimu
f1=[0]*nsimu
test_fraction =[0]*nsimu
for i in range(1,nsimu):
    X_train, X_test, y_train, y_test = train_test_split(dataframe.drop('Survived',axis=1),
dataframe['Survived'], test_size=0.1+(i-1)*0
random_state=111)          logmodel =(LogisticRegression(C=1,tol=1e-4, max_iter=1000,n_jobs=4))
logmodel.fit(X_train,y_train)          predictions = logmodel.predict(X_test)
    class_report[i] = classification_report(y_test,predictions)
l=class_report[i].split()          f1[i] = l[len(l)-2]
    test_fraction[i]=0.1+(i-1)*0.007

plt.plot(test_fraction[1:len(test_fraction)-2],f1[1:len(f1)-2])
plt.title("F1-score vs. test set size (fraction)",fontsize=20)
plt.xlabel("Test set size (fraction)",fontsize=17)
plt.ylabel("F1-score on test data",fontsize=17) plt.show()

```

F1-score as a function of random seed of test/train split

```

nsimu=101
class_report = [0]*nsimu
f1=[0]*nsimu
random_init
=[0]*nsimu
for i in
range(1,nsimu):
    X_train, X_test, y_train, y_test = train_test_split(dataframe.drop('Survived',axis=1),
dataframe['Survived'], test_size=0.3,
random_state=i+100)          logmodel =(LogisticRegression(C=1,tol=1e-5,
max_iter=1000,n_jobs=4))          logmodel.fit(X_train,y_train)          predictions =
logmodel.predict(X_test)
    class_report[i] = classification_report(y_test,predictions)
l=class_report[i].split()          f1[i] = l[len(l)-2]
random_init[i]=i+100

plt.plot(random_init[1:len(random_init)-2],f1[1:len(f1)-2])
plt.title("F1-score vs. random initialization seed",fontsize=20)
plt.xlabel("Random initialization seed",fontsize=17)
plt.ylabel("F1-score on test data",fontsize=17) plt.show()

```

Implement Kernel Density Estimation for Feature Space

Import required libraries

```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns from
scipy import stats
```

We take 100 random samples and find minimum and maximum in them.

`np.linspace(x_min, x_max, 100)`: This generates an array called `x_axis` that contains 100 evenly spaced values starting from `x_min` and ending at `x_max`. This is often used to create a range of x-values for plotting

```
dataset = np.random.randn(100)
x_min = dataset.min() - 2
x_max = dataset.max() + 2
x_axis = np.linspace(x_min, x_max, 100)
```

```
print(x_min, x_max)
```

```
-4.53534782641586  4.623552856673909
```

Calculating the bandwidth for kernel density estimation. The bandwidth determines the smoothness of the estimated probability density function (PDF) when using kernel density estimation. The formula you are using for bandwidth appears to be based on a practical estimation method.

The bandwidth calculated in this way will be used in kernel density estimation to control the width of the kernel, which affects the smoothness of the estimated PDF. The bandwidth can have a significant impact on the appearance and accuracy of kernel density plots.

```
#set up the bandwidth, for info on this: url =
'http://en.wikipedia.org/wiki/Kernel_density_estimation/#Practical_estimation_of_the_bandwidth'
bandwidth = ((4*dataset.std()**-0.5)/(3*len(dataset)))**0.2
bandwidth
0.42199405201453
```

```
#create an empty kernel list
kernel_list = []
```


```
#Plot each basis function
for data_point in dataset:
```

```
    #create a kernel for each point and append to list
    kernel = stats.norm(data_point, bandwidth).pdf(x_axis)
    kernel_list.append(kernel)
```

```
    #scale for plotting
    kernel = kernel / kernel.max()
    kernel = kernel * .4
    plt.plot(x_axis, kernel, color = 'violet', alpha = 0.5)
```

```
plt.ylim(0,1)
```

```
(0.0, 1.0)
1.0
0.8
0.6
0.4
0.2
0.0
#to get the kde plot we can sum these basis function,
#plot the sum of the basis function
sum_of_kde = np.sum(kernel_list,axis = 0)
#plot figure
fig = plt.plot(x_axis,sum_of_kde,color = "indianred")
#add the initial
rugplot
sns.rugplot(dataset,c = "indianred")
#get rid of y-tick marks
plt.yticks([])
#set title plt.suptitle("Sum of the
basis function")
Text(0.5, 0.98, 'Sum of the basis function')
```



Implement Dimensionality Reduction using Principal Component Analysis (PCA)

Get required libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns from sklearn
import preprocessing from
sklearn.decomposition import PCA
```

Rename the columns into desired names.

```
a = pd.read_csv("iris.data", names=['sepal length', 'sepal width', 'petal length', 'petal width', 'target'])
print(a)
a.shape
```

	sepal length	sepal width	petal length	petal width	target
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa..

145	6.7	3.0	5.2	2.3	Iris-virginica
146	6.3	2.5	5.0	1.9	Iris-virginica
147	6.5	3.0	5.2	2.0	Iris-virginica
148	6.2	3.4	5.4	2.3	Iris-virginica
149	5.9	3.0	5.1	1.8	Iris-virginica

150 rows x 5 columns] (150, 5)

```
features = a.columns
features
```

```
Index(['sepal length', 'sepal width', 'petal length', 'petal width', 'target'], dtype='object')
```

```
from sklearn.preprocessing import StandardScaler
features = ['sepal length', 'sepal width', 'petal length', 'petal width']
x = a.loc[:, features].values y = a.loc[:, ['target']].values x =
StandardScaler().fit_transform(x)
```

```
pca = PCA(n_components=2)
principalComponents = pca.fit_transform(x)
principalDataframe = pd.DataFrame(data = principalComponents, columns = ['PC1', 'PC2'])
```

```
targetDataframe = a[['target']]
newDataframe = pd.concat([principalDataframe, targetDataframe], axis = 1)
```

newDataframe

	PC1	PC2	target
0	-2.264542	0.505704	Iris-setosa
1	-2.086426	-0.655405	Iris-setosa
2	-2.367950	-0.318477	Iris-setosa

Creating a scatter plot of data points in a two-dimensional space defined by the Principal Component 1 (PC1) and Principal Component 2 (PC2). The plot is intended to visualize the relationship or distribution of data points in this reduced-dimensional space.

You can see how data points are distributed in the PC1-PC2 space. It helps you understand the concentration of data points and whether they...

... form clusters or exhibit patterns.

```
145 1.870522 0.382822 Iris-virginica
```

The scatter plot can show whether there is a correlation or relationship between PC1 and PC2. If the points tend to follow a linear trend or form

```
146 1.558492 -0.905314 Iris-virginica a distinct shape, it indicates a correlation between the two components. 147 1.520845 0.266795 Iris-virginica
```

Outliers, if present, may be visible as data points that deviate significantly from the general trend. These outliers can provide valuable

```
148 1.376391 1.016362 Iris-virginica
```

information about data anomalies.

```
149 0.959299 -0.022284 Iris-virginica
```


```
150 rows x 3 columns
```

```
plt.scatter(principalDataframe.PC1,principalDataframe.PC2)
```

```
plt.title("PC1 against PC2") plt.xlabel("PC1")
```

```
plt.ylabel("PC2")
```

```
Text(0, 0.5, 'PC2')
```



A scatter plot for data points in a two-dimensional space defined by Principal Component 1 (PC1) and Principal Component 2 (PC2). This plot appears to be specifically designed for visualizing the Iris dataset, which contains samples of three different species: Iris-setosa, Iris-versicolor, and Iris-virginica.

You will see a scatter plot with data points from the Iris dataset displayed in two dimensions (PC1 and PC2).

Data points belonging to different Iris species (setosa, versicolor, and virginica) will be distinguished by different colors (red, green, and blue).

The plot visually shows how the Iris species are distributed in the PC1-PC2 space. It can help you observe patterns, clusters, or separations between these species.

The legend will identify which color corresponds to each Iris species, making it easy to interpret the plot.

A scatter plot for data points in a two-dimensional space defined by Principal Component 1 (PC1) and Principal Component 2 (PC2). This plot appears to be specifically designed for visualizing the Iris dataset, which contains samples of three different species: Iris-setosa, Iris-versicolor, and Iris-virginica.

You will see a scatter plot with data points from the Iris dataset displayed in two dimensions (PC1 and PC2).

Data points belonging to different Iris species (setosa, versicolor, and virginica) will be distinguished by different colors (red, green, and blue).


The plot visually shows how the Iris species are distributed in the PC1-PC2 space. It can help you observe patterns, clusters, or separations between these species.

The legend will identify which color corresponds to each Iris species, making it easy to interpret the plot.

```

fig = plt.figure(figsize = (8,8))
ax = fig.add_subplot(1,1,1)
ax.set_xlabel("PC1")
ax.set_ylabel("PC2")
ax.set_title("Plot of PC1 vs PC2",fontsize = 20) targets =
['Iris-setosa','Iris-versicolor','Iris-virginica'] colors =
['r','g','b'] for target,color in zip(targets,colors):
    indicesToKeep = newDataframe['target'] == target
    ax.scatter(newDataframe.loc[indicesToKeep,'PC1'],newDataframe.loc[indicesToKeep,'PC2'],c = col
ax.legend(targets)  ax.grid()

```



```

explained_variance_ratio = pca.explained_variance_ratio_
explained_variance_ratio
array([0.72770452, 0.23030523])

```

Implement K-Means Clustering using Synthetic Data from

Problem:

You have a multidimensional set of data (such as a set of hidden unit activations) and you want to see which points are closest to others. The k-means algorithm searches for a pre-determined number of clusters within an unlabeled multidimensional dataset. It accomplishes this using a simple conception of what the optimal clustering looks like:

- The "cluster center" is the arithmetic mean of all the points belonging to the cluster.
- Each point is closer to its own cluster center than to other cluster centers.


```
import matplotlib.pyplot as plt
import seaborn as sns
sns.set()#plot styling import
numpy as np
```

Create synthetic dataset of unlabeled blobs

The dataset would be synthesized using sklearn.datasets.samples generator from the sklearn package. You will import binary large objectsblobs to form clusters from the synthetic dataset.

```
from sklearn.datasets import make_blobs
X, y_true = make_blobs(n_samples=300,centers = 4, cluster_std = 0.60, random_state = 0)
plt.scatter(X[:,0],X[:,1],s=50,color = 'blue')
```

<matplotlib.collections.PathCollection at 0x7f97c2206bf0>



Import K-means from Sklearn and Fit the data


Verify the syntentic dataset and t the data to the K-Means model.

```
from sklearn.cluster import KMeans kmeans =
KMeans(n_clusters = 4, n_init = 10)
kmeans.fit(X) y_kmeans = kmeans.predict(X)
```

Visualize the tted data by coloring the blobs aby assigned label numbers

Verify the syntentic dataset and t the data to the K-Means model.

```
plt.scatter(X[:,0],X[:,1],c = y_kmeans,s = 50, cmap = 'viridis')
centers = kmeans.cluster_centers_
plt.scatter(centers[:,0],centers[:,1],c='black',s = 200, alpha = .5);
```



How k-means is a special case of Expectation-maximization (EM) algorithm Expectation-maximization (EM) is a powerful algorithm that comes up in a variety of contexts within data science. k-means is a particularly simple and special case of this more general algorithm. The basic algorithmic flow of k-means is to

- Guess some cluster center (initialization)
 - Repeat following steps until converged
- E-step: assign points to the nearest cluster center
- M-Step: set the cluster centers to the mean

```
from sklearn.metrics import pairwise_distances_argmin
```

```
def find_clusters(X, n_clusters, rseed = 2):
```



```
    #1. Randomly choosed clusters    rng =
    np.random.RandomState(rseed)    i =
    rng.permutation(X.shape[0])[:n_clusters]
    centers = X[i]

    while True:
        #2a. Assign labels based on closest center
        labels = pairwise_distances_argmin(X, centers)
```

```
        #2b. Find new centers from means of points
        new_centers = np.array([X[labels == i].mean(0)
```

```
                                for i in range(n_clusters)])
    #2c. Check for convergence    if
    np.all(centers == new_centers):
        break
    centers = new_centers
    return centers, labels
```

```
centers, labels = find_clusters(X , 4)
plt.scatter(X[:,0], X[:,1], c = labels, s = 50, cmap = "viridis")
<matplotlib.collections.PathCollection at 0x7f97b1465510>
```




How many number of clusters

Plot the clusters formed using scatter plot

```
labels = KMeans(6, random_state = 0, n_init = 10).fit_predict(X)
plt.scatter(X[:, 0], X[:, 1], c = labels, s = 50, cmap = 'viridis')
```

<matplotlib.collections.PathCollection at 0x7f97b0ed5690>




Limitation of K-Means Algorithm

```
from sklearn.datasets import make_moons
from sklearn.datasets import make_moons
X, y = make_moons(200, noise = .05, random_state = 0)

labels = KMeans(2, random_state = 0, n_init = 10,).fit_predict(X)
plt.scatter(X[:, 0], X[:, 1], c = labels, s = 50, cmap = 'viridis')
```

<matplotlib.collections.PathCollection at 0x7f97b0f46500>



Kernel Transformation

The situation above is reminiscent of the Support Vector Machines, where we use a kernel transformation to project the data into a higher dimension where a linear separation is possible. We might imagine using the same trick to allow k-means to discover non-linear boundaries.


One version of this kernelized k-means is implemented in Scikit-Learn within the SpectralClustering estimator. It uses the graph of nearest neighbors to compute a higher- dimensional representation of the data, and then assigns labels using a k-means algorithm using the following code

```
from sklearn.cluster import SpectralClustering
model = SpectralClustering(n_clusters = 2, affinity = 'nearest_neighbors', assign_labels = 'kmeans')
labels = model.fit_predict(X)
plt.scatter(X[:, 0], X[:, 1], c = labels, s = 50, cmap = 'viridis')
```

/usr/local/lib/python3.10/dist-packages/sklearn/manifold/_spectral_embedding.py:274:

warnings.warn(

<matplotlib.collections.PathCollection at 0x7f97b0fe8be0>



Implement Gaussian Mixture Model using Synthetic Dataset

Challenges in K-Means can be overcome using:


- You could measure uncertainty in cluster assignment by comparing the distances of each point to all cluster centers, rather than focusing on just the closest.
- You might also imagine allowing the cluster boundaries to be ellipses rather than circles, so as to account for non-circular clusters.

```
import matplotlib.pyplot as plt
import seaborn as sns
sns.set()
#plot styling
import numpy as np
```

Generate Synthetic Data using unlabeled blobs

```
from sklearn.datasets import make_blobs
X, y_true = make_blobs(n_samples=400, centers=4,
cluster_std=0.7, random_state=0)
plt.scatter(X[:, 0], X[:, 1], s=50, color = 'blue')
```

<matplotlib.collections.PathCollection at 0x7bff50fbac80>



Generalize to Gaussian Mixture Models

```
from sklearn.mixture import GaussianMixture
gmm = GaussianMixture(n_components=8).fit(X)
labels = gmm.predict(X)
plt.scatter(X[:, 0], X[:, 1], c=labels, s = 40, cmap = 'viridis')
```

<matplotlib.collections.PathCollection at 0x7bff48212f20>



```

10
8
6
probs = gmm.predict_proba(X)
print(probs[: 5].round(3))

[[0.988 0.    0.006 0.    0.003 0.    0.    0.002]
 [0.    0.    0.992 0.    0.    0.003 0.    0.002]
 [0.    0.    0.994 0.    0.    0.006 0.    0.002]
 [0.746 0.    0.    0.253 0.    0.001 0.    0.    0.001]
 [0.    0.    0.989 0.    0.    0.011 0.    0.    0.]]

2
# print(probs.max(1))
size = probs.max(1)/0.03 # square emphasizes differences
#print(size)
plt.scatter(X[:, 0], X[:, 1], c = labels, edgecolor = 'k', cmap = 'viridis', s=size)

```




GMM as Density Estimation and Generative Model

```

from sklearn.datasets import make_moons
Xmoon, Ymoon = make_moons(200, noise = 0.08, random_state = 0)
plt.scatter(Xmoon[:, 0], Xmoon[:, 1], edgecolor = 'k')

```

<matplotlib.collections.PathCollection at 0x7bff434a2020>



The function `rst` ts the GMM to the data using the `t` method and then predicts cluster labels for each data point using the `predict` method. If `label` is `True`, it colors the data points based on their cluster labels using the 'viridis' colormap. It also plots ellipses representing the GMM's components using the `draw_ellipse` function, with the ellipses' properties being determined by the GMM's means, covariances, and weights.

The variable `w_factor` is used to adjust the transparency of the ellipses based on the weights of the GMM components. This makes the ellipses more transparent for components with lower weights.

The expected outcome when using these functions is a Matplotlib scatter plot where data points are colored according to their cluster assignments if '`label=True`'. In addition, ellipses will be drawn to represent the shape, orientation, and size of each Gaussian component in the GMM. This allows you to visually understand the clustering and characteristics of the data based on the GMM model.

```
from matplotlib.patches import Ellipse
def draw_ellipse(position, covariance, ax=None, **kwargs):
    """Draw an ellipse with a given position and covariance"""
    ax = ax or plt.gca()

    # Convert covariance to principal axes
    if covariance.shape == (2, 2):
        U, s, Vt = np.linalg.svd(covariance)
        angle = np.degrees(np.arctan2(U[1, 0], U[0, 0]))
        width, height = 2 * np.sqrt(s)
    else:
        angle = 0
        width, height = 2 * np.sqrt(covariance)

    # Draw the Ellipse
    for nsig in range(1, 4):
        ax.add_patch(Ellipse(position, nsig * width, nsig * height,
                              angle, **kwargs))

def plot_gmm(gmm, X, label=True, ax=None):
    ax = ax or plt.gca()
    labels = gmm.fit(X).predict(X)
    if label:
        ax.scatter(X[:, 0], X[:, 1], c=labels, s=40, cmap='viridis',
                  zorder=2, edgecolor='k')
    else:
        ax.scatter(X[:, 0], X[:, 1], s=40, zorder=2, cmap='viridis', edgecolor='k')
    ax.axis('equal')

    w_factor = 0.2 / gmm.weights_.max()
    for pos, covar, w in zip(gmm.means_, gmm.covariances_, gmm.weights_):
        draw_ellipse(pos, covar, alpha=w * w_factor)
```


The outcome of this code will be a Matplotlib plot where:

Data points from `Xmoon` will be scattered on the plot, and each point will be colored according to its cluster assignment.

Ellipses will be drawn around the clusters to represent the estimated shapes, orientations, and sizes of the GMM components.

```
gmm2 = GaussianMixture(n_components=4, covariance_type='full',
                        random_state=42)
plt.figure(figsize = (8,5))
plot_gmm(gmm2, Xmoon)
```

```
<ipython-input-71-18cf7d486bb7>:17: MatplotlibDeprecationWarning: Passing the angle parameter of __init__() positionally is deprecated
ax.add_patch(Ellipse(position, nsig * width, nsig * height,
```



Implement Support Vector Machine Classification using Breast Cancer Dataset

- ▼ In machine learning, support vector machines (SVMs, also support vector networks) are supervised learning models with associated learning algorithms that analyze data used for classification and regression analysis.

An SVM model is a representation of the examples as points in space, mapped so that the examples of the separate categories are divided by a clear gap that is as wide as possible. New examples are then mapped into that same space and predicted to belong to a category based on which side of the gap they fall. This gap is also called maximum margin and the SVM classifier is called maximum margin classifier.

In addition to performing linear classification, SVMs can efficiently perform a non-linear classification using what is called the kernel trick, implicitly mapping their inputs into high-dimensional feature spaces.

```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
```

Get the data



We'll use the built-in breast cancer dataset from Scikit-learn. Note the load function:

```
from sklearn.datasets import load_breast_cancer
cancer = load_breast_cancer()
```

The dataset is presented in a dictionary format

```
cancer.keys()

dict_keys(['data', 'target', 'frame', 'target_names', 'DESCR', 'feature_names', 'filename', 'data_module'])

cancer[
    'feature_names'
]

array(['mean radius', 'mean texture', 'mean perimeter', 'mean area',
       'mean smoothness', 'mean compactness', 'mean concavity',
       'mean concave points', 'mean symmetry', 'mean fractal dimension',
       'radius error', 'texture error', 'perimeter error', 'area error',
       'smoothness error', 'compactness error', 'concavity error',
       'concave points error', 'symmetry error',
       'fractal dimension error', 'worst radius', 'worst texture',
       'worst perimeter', 'worst area', 'worst smoothness',
       'worst compactness', 'worst concavity', 'worst concave points',
       'worst symmetry', 'worst fractal dimension'], dtype='<U23')
```

Set up dataframe



```
df = pd.DataFrame(cancer['data'], columns=cancer['feature_names'])
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 30 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   mean radius                           569 non-null    float64
1   mean texture                           569 non-null    float64
2   mean perimeter                         569 non-null    float64
3   mean area                             569 non-null    float64
4   mean smoothness                       569 non-null    float64
5   mean compactness                      569 non-null    float64
6   mean concavity                        569 non-null    float64
7   mean concave points                   569 non-null    float64
8   mean symmetry                         569 non-null    float64
9   mean fractal dimension                 569 non-null    float64
10  radius error                          569 non-null    float64
11  texture error                         569 non-null    float64
12  perimeter error                       569 non-null    float64
```

29/10/2023, 12:54

StatMLLab09.ipynb - Colaboratory

13

area error

569 non-null

float64

14

smoothness error

569 non-null

float64

15

compactness error

569 non-null

float64

16

concavity error

569 non-null

float64

17

concave points error

569 non-null

float64

18

symmetry error

569 non-null

float64

19

fractal dimension error

569 non-null

float64

20

worst radius

569 non-null

float64

21

worst texture

569 non-null

float64

22

worst perimeter

569 non-null

float64

23

worst area

569 non-null

float64

24

worst smoothness

569 non-null

float64

25

worst compactness

569 non-null

float64

26

worst concavity

569 non-null

float64

27

worst concave points

569 non-null

float64

28

worst symmetry

569 non-null

float64

29 worst fractal dimension 569 non-null float64 dtypes: float64(30)

memory usage: 133.5 KB

df.describe()

	mean radius			mean	mean	mean	mean	mean	mean	mean texture	mean area	mean perimeter	mean concave points	mean fractal dimension	...	
count	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	...	56
mean	14.127292	19.289649	91.969033	654.889104	0.096360	0.104341	0.088799	0.048919	0.181162	0.062798	1
std	3.524049	4.301036	24.298981	351.914129	0.014064	0.052813	0.079720	0.038803	0.027414	0.007060
min	6.981000	9.710000	43.790000	143.500000	0.052630	0.019380	0.000000	0.000000	0.106000	0.049960
25%	11.700000	16.170000	75.170000	420.300000	0.086370	0.064920	0.029560	0.020310	0.161900	0.057700	1
50%	13.370000	18.840000	86.240000	551.100000	0.095870	0.092630	0.061540	0.033500	0.179200	0.061540	1
75%	15.780000	21.800000	104.100000	782.700000	0.105300	0.130400	0.130700	0.074000	0.195700	0.066120	1
max	28.110000	39.280000	188.500000	2501.000000	0.163400	0.345400	0.426800	0.201200	0.304000	0.097440	3
8 rows × 30 columns																

np.sum(pd.isnull(df).sum()) #Sum of the count of the null objects in all

0

What are the 'target' data in the dataset?

cancer['target'].sum()

357

Adding the target data to DataFrame

df['Cancer'] = pd.DataFrame(cancer['target'])

df.head()


	mean							mean		mean		mean		mean	
	radius	texture	perimeter	area	smoothness	worst texture	worst perimeter	compactness	worst concave points	concavity	fractal dimension	...	symmetry	...	mean
0	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	0.2419	0.07871	...	17.33	184.60	2019	
1	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	0.1812	0.05667	...	23.41	158.80	1956	
2	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	0.2069	0.05999	...	25.53	152.50	1709	
3	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	0.2597	0.09744	...	26.50	98.87	567	
4	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	0.1809	0.05883	...	16.67	152.20	1575	
5	rows × 31 columns														

Exploratory Data analysis

**Check the relative counts of benign (0) vs malignant (1) cases of **





```
sns.set_style('whitegrid')
sns.countplot(x='Cancer', data=df, palette='RdBu_r')
```


<Axes: xlabel='Cancer', ylabel='count'>




Draw boxplots of all the mean features(rst 10 columns) for '0' and '1' CANCER OF

```
l = list(df.columns[0:10])
for i in range(len(l)-1):
    sns.boxplot(x='Cancer', y=l[i], data=df, palette='winter')
plt.figure()
```







<Figure size 640x480 with 0 Axes>

▼ Training and Prediction

Train Test Split

```
df_feat = df.drop('Cancer', axis=1)
df_feat.head()
```

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean worst compactness	mean concave points	mean concave dimension	mean fractal symmetry	mean radius		
0	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	0.2419	0.07871	25.38	17.33	184
1	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	0.1812	0.05667	24.99	23.41	158
2	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	0.2069	0.05999	23.57	25.53	152
3	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	0.2597	0.09744	14.91	26.50	98
4	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	0.1809	0.05883	22.54	16.67	152
5	rows × 30 columns												

```
df_target = df['Cancer']
df_target.head()

0    0
1    0
2    0
3    0
4    0 Name: Cancer, dtype: int64
```

```
from sklearn.model_selection import train_test_split
```



```
X_train, X_test, Y_train, Y_test = train_test_split(df_feat, df_target, test_size = .30,
                                                    random_state=101)

Y_train.head()

178    1
421    1
57     0
514    0
548    1
Name: Cancer, dtype: int64
```

Train the Support Vector Classifier

```
from sklearn.svm import SVC

model = SVC()

model.fit(X_train, Y_train)
```

```

SVC
SVC()
```

Predictions and Evaluations

```
predictions = model.predict(X_test)

from sklearn.metrics import
classification_report, confusion_matrix
```

Notice that we are classifying everything into a single class! This means our model need to normalize the data

```
print(confusion_matrix(Y_test, predictions))

[[ 56 10]
 [ 3 102]]
```

As expected the classification report card is bad

```
print(classification_report(Y_test, predictions))
```

		precision	recall	f1-score	support
	0	0.95	0.85	0.90	66
1	0.91	0.97	0.94	105	
	accuracy			0.92	171
	macro avg	0.93	0.91	0.92	171
	weighted avg	0.93	0.92	0.92	171

```
param_grid = {
    'C': [0.1, 1, 10, 100, 1000], 'gamma': [1, 0.1, 0.01, 0.001, 0.0001], 'kernel': ['rbf']
}

from sklearn.model_selection import GridSearchCV
grid =
```

```
GridSearchCV(SVC(), param_grid, refit=True, verbose=1)
```

```
#May take awhile
grid.fit(X_train, Y_train)
Fitting 5 folds for each of 25 candidates, totalling 125 fits
```

```

GridSearchCV
grid.best_estimator_
{'C': 1, 'gamma': 0.0001, 'kernel': 'rbf'}
grid.best_estimator_
```

▼ SVC

SVC(C=1, gamma=0.0001)

```
grid_predictions = grid.predict(X_test)
```

```
print(confusion_matrix(Y_test, grid_predictions))
```

```
[[ 59   7]
 [   4 101]]
```

```
print(classification_report(Y_test, grid_predictions))
```

		precision	recall	f1-score	support
	0	0.94	0.89	0.91	66
1	0.94	0.96	0.95	105	
accuracy				0.94	171
macro avg	0.94	0.93	0.93	171	
weighted avg	0.94	0.94	0.94	171	